

Film Industry Investment Analysis



Film Industry Vectors by Vecteezy (<https://www.vecteezy.com/free-vector/film-industry>)

Overview

This project analyzes which types of movies generate the most revenue at the box office. Descriptive analysis of movie data sets collected from Box Office Mojo, IMBD, Rotten Tomatoes, The Movie DB, and The Numbers, will highlight specific movie characteristics that increase ROI such as, genre selection, month of the film's release, and the length/runtime of the movie. This analysis can be used by the company to make decisions on these specific movie criteria and help the company's new movie studio operate with the lowest risk in their new business endeavor.

Business Problem



Film Industry Vectors by Vecteezy (<https://www.vecteezy.com/free-vector/film-industry>)

Analyzing several movie datasets can provide valuable insights that can lead to concrete business recommendations in various aspects of the film industry. When considering characteristics of a successful movie, the movie data was analyzed to identify movies with the greatest return on investment. Based on this analysis, which is explained below, are the following three recommendations:

1. Action and/or Adventure genre movies gross higher than other genres**
2. Most successful release months are May, June, November, and December**
3. The most profitable movies are 85-125 minutes in length**

It's important to note that any business recommendations derived from data analysis should be accompanied by careful consideration of the specific context, legal and regulatory requirements, and limitations of the dataset. Additionally, these recommendations should be subject to ongoing evaluation and refinement based on updated data and emerging industry practices.

Data Understanding

Some media databases provide public data that includes various data about historical and current films both domestically and internationally. Each database has a variety of variables some of which may match other databases. The quantity of data provided was excellent, however, the quality of these various datasets were weak.

Some files are compressed CSV (comma-separated values) or TSV (tab-separated values) files that can be opened using spreadsheet software or pd.read_csv, while the data from IMDb is located in a SQLite database.

I have been provided a diagram for the IMDb data and will have to explore the other data to see how they connect to each other and the IMDb data.

Overview of dataset:

1. Independent Sources:

- Internet Movie Database (IMDb)
- The Movie Database (TMDb)
- Box Office Mojo
- The Numbers

2. About: Selected film industry data reported between 1915-2020. Due to the differences in the independent databases, when choosing variables several considerations were made:

- Data Relevance
- Data Quantity
- Data Quality

3. Decision-based Variables used:

- Profit/Profit Margin
- Average Domestic Gross
- Movie Runtime
- Movie Release Date
- Genre of Movie

4. Missing Values: In the case of data quality, missing data were handled on a case by case basis. Some considerations made were the relative importance of the variable in a particular analysis, how much data was needed, and finally, the possibility of filling data with a central measure of tendency and the statistical implications.

Data Exploration

```
In [1]: 1 # Import the necessary packages.  
2  
3 # Native Python Packages  
4 import os  
5 from collections import Counter  
6 from glob import glob  
7 import functools  
8  
9 # Third Party Packages  
10 import pandas as pd  
11 import numpy as np  
12 import matplotlib.pyplot as plt  
13 from matplotlib.ticker import FuncFormatter  
14 import plotly.graph_objects as go  
15 import plotly.express as px  
16 import seaborn as sns  
17 import locale  
18 import inflect  
19 import sqlite3  
20 import cpi  
21 %matplotlib inline  
22
```

→ There are a number of .csv and .tsv files provided in the zipped folder, glob will be used to read them in by matching path names.

```
In [2]: 1 # Create a list of all csv files using glob  
2 data_folder = './Data'  
3 csv_files = glob(os.path.join(data_folder, '*.csv.gz'))  
4 for i, file in enumerate(csv_files, 1):  
5     print(i, file, sep=': ')  
  
1: ./Data/tmdb.movies.csv.gz  
2: ./Data/tn.movie_budgets.csv.gz  
3: ./Data/bom.movie_gross.csv.gz
```

→ Using the Name of the file as the key and and the Pandas Dataframe as the value, dictionaries will be used to store dataframes. All data can be previewed by iterrating over the dictionary this way.

In [3]:

```
1 # Create a dictionary of DataFrames
2 csv_dict = {}
3 for file in csv_files:
4     filename = os.path.splitext(os.path.basename(file))[0].replace('.csv', '').replace('.', '_') # Clean file name
5     file_df = pd.read_csv(file) # Create DataFrame
6     csv_dict[filename] = file_df # Insert DataFrame into dictionary
```

In [4]:

```

1 # Preview the data
2 for name, df in csv_dict.items():
3     print(name)
4     print('Total number of results:', len(df))
5     display(df.head()) # Using display instead of print leads to neater formatting in Jupyter Notebook

```

tmdb_movies

Total number of results: 26517

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | vote_average | vote_count |
|---|---------------|------------------------|-------|-------------------|--|------------|--------------|--|--------------|------------|
| 0 | 0 | [12, 14, 10751] | 12444 | en | Harry Potter and the Deathly Hallows: Part 1 | 33.533 | 2010-11-19 | Harry Potter and the Deathly Hallows: Part 1 | 7.7 | 10788 |
| 1 | 1 | [14, 12, 16, 10751] | 10191 | en | How to Train Your Dragon | 28.734 | 2010-03-26 | How to Train Your Dragon | 7.7 | 7610 |
| 2 | 2 | [12, 28, 878] | 10138 | en | Iron Man 2 | 28.515 | 2010-05-07 | Iron Man 2 | 6.8 | 12368 |
| 3 | 3 | [16, 35, 10751] | 862 | en | Toy Story | 28.005 | 1995-11-22 | Toy Story | 7.9 | 10174 |
| 4 | 4 | [28, 878, 12] | 27205 | en | Inception | 27.920 | 2010-07-16 | Inception | 8.3 | 22186 |

tn_movie_budgets

Total number of results: 5782

| | id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|---|----|--------------|---|-------------------|----------------|-----------------|
| 0 | 1 | Dec 18, 2009 | Avatar | \$425,000,000 | \$760,507,625 | \$2,776,345,279 |
| 1 | 2 | May 20, 2011 | Pirates of the Caribbean: On Stranger Tides | \$410,600,000 | \$241,063,875 | \$1,045,663,875 |
| 2 | 3 | Jun 7, 2019 | Dark Phoenix | \$350,000,000 | \$42,762,350 | \$149,762,350 |
| 3 | 4 | May 1, 2015 | Avengers: Age of Ultron | \$330,600,000 | \$459,005,868 | \$1,403,013,963 |
| 4 | 5 | Dec 15, 2017 | Star Wars Ep. VIII: The Last Jedi | \$317,000,000 | \$620,181,382 | \$1,316,721,747 |

bom_movie_gross

Total number of results: 3387

| | title | studio | domestic_gross | foreign_gross | year |
|---|---|--------|----------------|---------------|------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000 | 2010 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000 | 2010 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000 | 2010 |
| 3 | Inception | WB | 292600000.0 | 535700000 | 2010 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000 | 2010 |

→ Again, use the same process for the .tsv files.

```
In [5]: 1 # Create a list of all tsv files using glob
2 data_folder = './Data'
3 tsv_files = glob(os.path.join(data_folder, '*.tsv.gz'))
```

```
In [6]: 1 # Create a dictionary of DataFrames
2 tsv_dict = {}
3 for file in tsv_files:
4     filename = os.path.splitext(os.path.basename(file))[0].replace('.tsv', '').replace('.', '_') # Clean
5     file_df = pd.read_table(file, delimiter = '\t', encoding = 'latin-1') # Create DataFrame
6     tsv_dict[filename] = file_df # Insert DataFrame into dictionary
```

In [7]:

```
1 # Preview the data
2 for name, df in tsv_dict.items():
3     print(name)
4     print('Total number of results:', len(df))
5     display(df.head()) # Using display instead of print leads to neater formatting in Jupyter Notebook
```

rt_reviews

Total number of results: 54432

| | id | review | rating | fresh | critic | top_critic | publisher | date |
|---|-----------|---|---------------|--------------|----------------|-------------------|------------------|-------------------|
| 0 | 3 | A distinctly gallows take on contemporary fina... | 3/5 | fresh | PJ Nabarro | 0 | Patrick Nabarro | November 10, 2018 |
| 1 | 3 | It's an allegory in search of a meaning that n... | NaN | rotten | Annalee Newitz | 0 | io9.com | May 23, 2018 |
| 2 | 3 | ... life lived in a bubble in financial dealin... | NaN | fresh | Sean Axmacher | 0 | Stream on Demand | January 4, 2018 |
| 3 | 3 | Continuing along a line introduced in last yea... | NaN | fresh | Daniel Kasman | 0 | MUBI | November 16, 2017 |
| 4 | 3 | ... a perverse twist on neorealism... | NaN | fresh | NaN | 0 | Cinema Scope | October 12, 2017 |

rt_movie_info

Total number of results: 1560

| | id | synopsis | rating | genre | director | writer | theater_date | dvd_date | currency | box_office | runtime | status |
|---|-----------|--|---------------|--|--|---|---------------------------------|----------------------------------|-----------------|--------------------|--------------------------------|---------------|
| 0 | 1 | This gritty, fast-paced, and innovative police... New York City, not-too-distant-future: Eric Pa... | R | Action and Adventure Classics Drama Drama Science Fiction and Fantasy | William Friedkin David Cronenberg | Ernest Tidyman David Cronenberg Don DeLillo | Oct 9, 1971 Aug 17, 2012 | Sep 25, 2001 Jan 1, 2013 | NaN \$ | NaN 600,000 | 104 minutes 108 minutes | Entertain... |
| 2 | 5 | Illeana Douglas delivers a superb performance ... Michael Douglas runs afoul of a treacherous su... | R | Drama Musical and Performing Arts Drama Mystery and Suspense | Allison Anders Barry Levinson | Allison Anders Paul Attanasio Michael Crichton | Sep 13, 1996 Dec 9, 1994 | Apr 18, 2000 Aug 27, 1997 | NaN NaN | NaN NaN | 116 minutes 128 minutes | Im... |
| 4 | 7 | NaN | NR | Drama Romance | Rodney Bennett | Giles Cooper | NaN | NaN | NaN | NaN | 200 minutes | Im... |

→ There are a number of DataFrames. This analysis will focus on the DataFrames that likely lead to recommendations.

The first DataFrame that stands out is tn_movie_budgets. This table shows the production budget, domestic and worldwide gross. This financial information will be valuable in exploring the meaning of profit. It can also be used as a standard measure of success based on which films performed best in the box office.

Additional tables that will be explored include: bom_movie_gross (domestic and worldwide profit), and tmdb_movies (genre). movie_basics (genre, runtime) and will also be incorporating data from an SQL database.

Data Preparation

Data Cleaning

The request from the investors is to recommend movie characteristics that maximize their return on investment. Addressing missing values across these databases requires thoughtful consideration, as relevant data is presented in different databases and so there must be clean and standardized variables in order to connect the data across databases.

In [8]:

```
1 # Clean up 'bom_movie_gross'  
2 movie_gross = csv_dict['bom_movie_gross']  
3 # Get summary  
4 movie_gross.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3387 entries, 0 to 3386  
Data columns (total 5 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   title            3387 non-null    object    
 1   studio           3382 non-null    object    
 2   domestic_gross   3359 non-null    float64  
 3   foreign_gross    2037 non-null    object    
 4   year             3387 non-null    int64    
dtypes: float64(1), int64(1), object(3)  
memory usage: 132.4+ KB
```

→ movie_gross will require that we convert the money related data types to integers.

In [9]:

```
1  """
2      calculates the .... what (not how) - minimal impletation details, if needed.
3  -
4  Input:
5  df : Pandas dataframe, include the relavent column names
6  -
7  Output:
8  summary : Pandas dataframe, details about new dataframe
9
10 #####in line comments are what and how#####(the one liners above the code)
11 """
12
13 # Define function to clean numbers
14 def numclean(df, col):
15     """
16     Formats the float64 and the object values into integers. All
17     NaN values within a specified column will be dropped from the data frame.
18     -
19     Input:
20     df : Pandas dataframe, include the relavent column names
21     col: A string that represents a column name
22     -
23     Output:
24     df : Pandas dataframe, with integer and without Nan values
25     """
26
27     # Convert to numeric, handling NaN values
28     df[col] = pd.to_numeric(df[col], errors='coerce', downcast='integer')
29
30     # Drop rows with NaN values in the specified column
31     df = df.dropna(subset=[col])
32
33     return df
34
```

```
In [10]: 1 # Run numclean function on the two columns
2 numclean(movie_gross,'foreign_gross')
3 numclean(movie_gross,'domestic_gross')
4
5 # Display dataframe with both columns cleaned
6 display(movie_gross.head())
```

| | | title | studio | domestic_gross | foreign_gross | year |
|---|--|---|--------|----------------|---------------|------|
| 0 | | Toy Story 3 | BV | 415000000.0 | 652000000.0 | 2010 |
| 1 | | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000.0 | 2010 |
| 2 | | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000.0 | 2010 |
| 3 | | Inception | WB | 292600000.0 | 535700000.0 | 2010 |
| 4 | | Shrek Forever After | P/DW | 238700000.0 | 513900000.0 | 2010 |

```
In [11]: 1 # Overview the data for an potential anomalies
2 movie_gross.describe()
```

Out[11]:

| | domestic_gross | foreign_gross | year |
|-------|----------------|---------------|-------------|
| count | 3.359000e+03 | 2.032000e+03 | 3387.000000 |
| mean | 2.874585e+07 | 7.505704e+07 | 2013.958075 |
| std | 6.698250e+07 | 1.375294e+08 | 2.478141 |
| min | 1.000000e+02 | 6.000000e+02 | 2010.000000 |
| 25% | 1.200000e+05 | 3.775000e+06 | 2012.000000 |
| 50% | 1.400000e+06 | 1.890000e+07 | 2014.000000 |
| 75% | 2.790000e+07 | 7.505000e+07 | 2016.000000 |
| max | 9.367000e+08 | 9.605000e+08 | 2018.000000 |

→ There aren't any entries here that seem unusual. It's worth noting that the movies in this data base range in year from 2010-2018.

tn_movie_budgets will require that the:

1. '\$' symbol

2. the comma separator be removed
3. the number converted to an integer

```
In [12]: 1 #Clean up 'tn_movie_gross'  
2 ww_gross = csv_dict['tn_movie_budgets']  
3 display(ww_gross.head())
```

| | id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|---|-----------|---------------------|---|--------------------------|-----------------------|------------------------|
| 0 | 1 | Dec 18, 2009 | Avatar | \$425,000,000 | \$760,507,625 | \$2,776,345,279 |
| 1 | 2 | May 20, 2011 | Pirates of the Caribbean: On Stranger Tides | \$410,600,000 | \$241,063,875 | \$1,045,663,875 |
| 2 | 3 | Jun 7, 2019 | Dark Phoenix | \$350,000,000 | \$42,762,350 | \$149,762,350 |
| 3 | 4 | May 1, 2015 | Avengers: Age of Ultron | \$330,600,000 | \$459,005,868 | \$1,403,013,963 |
| 4 | 5 | Dec 15, 2017 | Star Wars Ep. VIII: The Last Jedi | \$317,000,000 | \$620,181,382 | \$1,316,721,747 |

```
In [13]: 1 # Convert release_date to a datetime object  
2 ww_gross['release_date'] = pd.to_datetime(ww_gross['release_date'])
```

```
In [14]: 1 # Create year column  
2 ww_gross['year'] = pd.DatetimeIndex(ww_gross['release_date']).year
```

In [15]:

```

1 # Define function to clean numbers
2 def numclean2(df,col):
3     """
4         Formats the float64 and the object values into integers. All
5         NaN values within a specified column will be dropped from the data frame.
6         Removes $ symbol, removes ',' separator and changes to int64 type
7     -
8     Input:
9     df : Pandas dataframe, include the relevant column names
10    col: A string that represents a column name
11    -
12    Output:
13    df : Pandas dataframe, with integer and without Nan values
14    """
15
16    if df[col].dtype != 'int64':
17        df[col]=df[col].str.replace('$','').str.replace(',','').astype('int64')
18
19    return df

```

In [16]:

```

1 # Run numclean2 function on the two columns
2 numclean2(ww_gross,'worldwide_gross')
3 numclean2(ww_gross,'domestic_gross')
4 numclean2(ww_gross,'production_budget')
5
6 # Display dataframe with all three columns cleaned
7 display(ww_gross.head())

```

| | id | release_date | movie | production_budget | domestic_gross | worldwide_gross | year |
|----------|-----------|---------------------|---|--------------------------|-----------------------|------------------------|-------------|
| 0 | 1 | 2009-12-18 | Avatar | 425000000 | 760507625 | 2776345279 | 2009 |
| 1 | 2 | 2011-05-20 | Pirates of the Caribbean: On Stranger Tides | 410600000 | 241063875 | 1045663875 | 2011 |
| 2 | 3 | 2019-06-07 | Dark Phoenix | 350000000 | 42762350 | 149762350 | 2019 |
| 3 | 4 | 2015-05-01 | Avengers: Age of Ultron | 330600000 | 459005868 | 1403013963 | 2015 |
| 4 | 5 | 2017-12-15 | Star Wars Ep. VIII: The Last Jedi | 317000000 | 620181382 | 1316721747 | 2017 |

```
In [17]: 1 # Overview the data for an anomalies  
2 ww_gross.describe()
```

Out[17]:

| | id | production_budget | domestic_gross | worldwide_gross | year |
|--------------|-------------|--------------------------|-----------------------|------------------------|-------------|
| count | 5782.000000 | 5.782000e+03 | 5.782000e+03 | 5.782000e+03 | 5782.000000 |
| mean | 50.372363 | 3.158776e+07 | 4.187333e+07 | 9.148746e+07 | 2003.967139 |
| std | 28.821076 | 4.181208e+07 | 6.824060e+07 | 1.747200e+08 | 12.724386 |
| min | 1.000000 | 1.100000e+03 | 0.000000e+00 | 0.000000e+00 | 1915.000000 |
| 25% | 25.000000 | 5.000000e+06 | 1.429534e+06 | 4.125415e+06 | 2000.000000 |
| 50% | 50.000000 | 1.700000e+07 | 1.722594e+07 | 2.798445e+07 | 2007.000000 |
| 75% | 75.000000 | 4.000000e+07 | 5.234866e+07 | 9.764584e+07 | 2012.000000 |
| max | 100.000000 | 4.250000e+08 | 9.366622e+08 | 2.776345e+09 | 2020.000000 |

→ This data table shows entries where the movie gross is zero and that's unusual. Some movies may have staggered release dates, may not be released in the USA, or simply the data was never entered and so there may not be values available for every movie title.

I will join these two data sets and then proceed to drop any data that still contains a zero value for the gross revenue.

Additionally, the range of years represented in this data table span from 1915-2020.

Profit and Profit Margin

The value of this venture relies on making a profit on the investment. Due to multiple revenue streams, it's challenging to understand whether or not a box office movie is ultimately profitable - and to whom. There is an abundance of information that takes careful consideration, in context, to interpret meaningfully.

For both a precise and accurate analysis, a concrete understanding of which variables the company would be owning would be needed. Due to the lack of detailed financial data available publically around movie profitability, this project will share a simplified analysis model based on publically available data.

In this analysis:

domestic profit is defined as follows: **profit = domestic gross revenue - production budget**

profit margin is defined to be: **profit margin = (domestic gross revenue - production budget) / domestic gross revenue *100**

*all values in this analysis have been adjusted for inflation for accurate comparison

Aggregating money data

```
In [18]: 1 # Create variables for dictionaries
2 rt_movie_info = tsv_dict['rt_movie_info']
3 tmdb_movies = csv_dict['tmdb_movies']
```

```
In [19]: 1 # Print columns for tmdb_movies data frame
2 print('Columns in tmdb_movies:')
3 print(tmdb_movies.columns)
4
5 # Print columns for tn_movie_budgets data frame
6 print('Columns in ww_gross:')
7 print(ww_gross.columns)
8
9 # Print columns for bom_movie_gross data frame
10 print('Columns in movie_gross:')
11 print(movie_gross.columns)
```

```
Columns in tmdb_movies:
Index(['Unnamed: 0', 'genre_ids', 'id', 'original_language', 'original_title',
       'popularity', 'release_date', 'title', 'vote_average', 'vote_count'],
      dtype='object')
Columns in ww_gross:
Index(['id', 'release_date', 'movie', 'production_budget', 'domestic_gross',
       'worldwide_gross', 'year'],
      dtype='object')
Columns in movie_gross:
Index(['title', 'studio', 'domestic_gross', 'foreign_gross', 'year'], dtype='object')
```

```
In [20]: 1 # Rename columns in ww_gross DataFrame
2 ww_gross = ww_gross.rename(columns={'movie': 'title'})
```

```
In [21]: 1 # How many entries of domestic gross are zero or null?
2 len(movie_gross[(movie_gross['domestic_gross'] == 0) | (movie_gross['domestic_gross'].isnull())])
```

```
Out[21]: 28
```

```
In [22]: 1 # How many entries of domestic gross are zero or null?  
2 len(ww_gross[(ww_gross['domestic_gross'] == 0) | (ww_gross['domestic_gross'].isnull())])
```

Out[22]: 548

```
In [23]: 1 # How many entries of domestic gross are zero or null?  
2 len(ww_gross[(ww_gross['worldwide_gross'] == 0) | (ww_gross['worldwide_gross'].isnull())])
```

Out[23]: 367

```
In [24]: 1 # How many entries of foreign gross are zero or null?  
2 len(movie_gross[(movie_gross['foreign_gross'] == 0) | (movie_gross['foreign_gross'].isnull())])
```

Out[24]: 1355

```
In [25]: 1 # How many entries of production budget are zero or null?  
2 len(ww_gross[(ww_gross['production_budget'] == 0) | (ww_gross['production_budget'].isnull())])
```

Out[25]: 0

→ Given the number of null or zero values between the movie_gross and ww_gross data sets, it turns out that the total number of usable data points reporting a worldwide gross is 5415 and the total number of usable data points reporting a domestic gross is 8565.

The majority of the foreign gross reports are zero or null values so this data will not be used.

Meta data could be used to verify whether all of the worldwide gross reports are in USD.

It makes sense to maximize the data points reporting revenue in order to advise stakeholders on the best way to increase ROI. All subsequent calculations will be based on available adjusted domestic revenue figures.

```
In [26]: 1 # Convert all movie title strings to lowercase  
2 movie_gross['title'] = movie_gross['title'].str.lower()  
3 ww_gross['title'] = ww_gross['title'].str.lower()
```

```
In [27]: 1 # Join the DataFrames using an 'outer' join to include all rows from both DataFrames
2 money_df = pd.merge(movie_gross, ww_gross, on=['title', 'year'], how='outer')
3 print('Total number of results:', len(money_df))
4 display(money_df.head())
```

Total number of results: 7914

| | title | studio | domestic_gross_x | foreign_gross | year | id | release_date | production_budget | domestic_gross_y | worldwide_gross |
|---|---|--------|------------------|---------------|------|------|--------------|-------------------|------------------|-----------------|
| 0 | toy story 3 | BV | 415000000.0 | 652000000.0 | 2010 | 47.0 | 2010-06-18 | 200000000.0 | 415004880.0 | 1.068880e+09 |
| 1 | alice in wonderland (2010) | BV | 334200000.0 | 691300000.0 | 2010 | NaN | NaT | NaN | NaN | NaN |
| 2 | harry potter and the deathly hallows part 1 | WB | 296000000.0 | 664300000.0 | 2010 | NaN | NaT | NaN | NaN | NaN |
| 3 | inception | WB | 292600000.0 | 535700000.0 | 2010 | 38.0 | 2010-07-16 | 160000000.0 | 292576195.0 | 8.355246e+08 |
| 4 | shrek forever after | P/DW | 238700000.0 | 513900000.0 | 2010 | 27.0 | 2010-05-21 | 165000000.0 | 238736787.0 | 7.562447e+08 |

```
In [28]: 1 # Convert release_date to a datetime object
2 money_df['release_date'] = pd.to_datetime(money_df['release_date'])
```

In [29]:

```
1 # Create columns and average values from two different data sets
2
3 # Calculate the average for domestic gross
4 money_df['avg_domestic_gross'] = money_df[['domestic_gross_x', 'domestic_gross_y']].replace(0, np.nan)
5
6 # Drop unnecessary columns
7 money_df.drop(['domestic_gross_x', 'domestic_gross_y', 'studio', 'id'], axis=1, inplace=True)
8
9 # Drop rows with NaN in 'avg_domestic_gross'
10 money_df.dropna(subset=['avg_domestic_gross'], inplace=True)
11
12 print('Total number of results:', len(money_df))
13 display(money_df.head())
```

Total number of results: 7344

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross |
|---|--|---|---------------|------|--------------|-------------------|-----------------|--------------------|
| 0 | | toy story 3 | 652000000.0 | 2010 | 2010-06-18 | 200000000.0 | 1.068880e+09 | 415002440.0 |
| 1 | | alice in wonderland (2010) | 691300000.0 | 2010 | NaT | NaN | NaN | 334200000.0 |
| 2 | | harry potter and the deathly hallows part 1 | 664300000.0 | 2010 | NaT | NaN | NaN | 296000000.0 |
| 3 | | inception | 535700000.0 | 2010 | 2010-07-16 | 160000000.0 | 8.355246e+08 | 292588097.5 |
| 4 | | shrek forever after | 513900000.0 | 2010 | 2010-05-21 | 165000000.0 | 7.562447e+08 | 238718393.5 |

In [30]:

```
1 # Accessing sorted duplicated movie titles
2 print('Total number of results:', len(money_df[money_df.duplicated(['title'], keep=False)]))
3 money_df[money_df.duplicated(['title'], keep=False)].sort_values(by=['title']).head()
4
```

Total number of results: 191

Out[30]:

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross |
|------|------------------------------|------------|---------------|------|--------------|-------------------|-----------------|--------------------|
| 7752 | 20,000 leagues under the sea | | NaN | 1916 | 1916-12-24 | 200000.0 | 8000000.0 | 8000000.0 |
| 6574 | 20,000 leagues under the sea | | NaN | 1954 | 1954-12-23 | 5000000.0 | 28200000.0 | 28200000.0 |
| 4315 | a monster calls | | NaN | 2017 | 2017-01-06 | 43000000.0 | 46414530.0 | 3740823.0 |
| 2435 | a monster calls | 43600000.0 | 2016 | | NaT | NaN | NaN | 3700000.0 |
| 4531 | a nightmare on elm street | | NaN | 2010 | 2010-04-30 | 35000000.0 | 117729621.0 | 63075011.0 |

In [31]:

```
1 # Sort the DataFrame by 'title' and 'release_date'
2 money_df.sort_values(by=['title', 'release_date'], inplace=True)
3
4 # Find duplicates within one year of each other for each title
5 mask = money_df.duplicated(subset=['title'], keep=False)
6 duplicates = money_df[mask]
7
8 # Create a mask to identify duplicates with a difference of 1 in years
9 mask_within_one_year = (duplicates['year'].diff().abs() == 1)
10
11 # Filter out duplicates where the difference between years is 1
12 duplicates_within_one_year = duplicates[mask_within_one_year]
13
14 # Keep the rows with the 'release_date' and drop the other duplicates
15 to_drop = duplicates_within_one_year[duplicates_within_one_year['release_date'].isnull()].index
16 money_df.drop(to_drop, inplace=True)
17
18 print('Total number of results:', len(money_df))
19 display(money_df.head())
```

Total number of results: 7333

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross |
|------|------------------------|-------|---------------|------|--------------|-------------------|-----------------|--------------------|
| 2093 | | '71 | 355000.0 | 2015 | NaT | NaN | NaN | 1300000.0 |
| 6309 | (500) days of summer | | NaN | 2009 | 2009-07-17 | 7500000.0 | 34439060.0 | 32425665.0 |
| 1830 | 1,000 times good night | | NaN | 2014 | NaT | NaN | NaN | 53900.0 |
| 2395 | 10 cloverfield lane | | 38100000.0 | 2016 | 2016-03-11 | 5000000.0 | 108286422.0 | 72091499.5 |
| 5913 | 10 days in a madhouse | | NaN | 2015 | 2015-11-11 | 12000000.0 | 14616.0 | 14616.0 |

In [32]: 1 money_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7333 entries, 2093 to 5331
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title            7333 non-null    object  
 1   foreign_gross     1994 non-null    float64 
 2   year              7333 non-null    int64   
 3   release_date      5238 non-null    datetime64[ns]
 4   production_budget 5238 non-null    float64 
 5   worldwide_gross   5238 non-null    float64 
 6   avg_domestic_gross 7333 non-null    float64 
dtypes: datetime64[ns](1), float64(4), int64(1), object(1)
memory usage: 458.3+ KB
```

In [33]:

```

1 # Adjust $ columns for inflation using CPI
2 columns_to_adjust = ['worldwide_gross', 'avg Domestic_gross', 'foreign_gross', 'production_budget']
3
4 for column in columns_to_adjust:
5     print(column)
6     money_df[f'cpi_{column}'] = money_df.apply(lambda row: cpi.inflate(row[column], row['year']), axis=1)
7
8 print('Total number of results:', len(money_df))
9 display(money_df.head())

```

worldwide_gross
 avg Domestic_gross
 foreign_gross
 production_budget
 Total number of results: 7333

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg Domestic_gross | cpi_worldwide_gross | cpi_avg Domestic_gross |
|------|------------------------|-------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|------------------------|
| 2093 | | '71 | 355000.0 | 2015 | NaT | NaN | NaN | 1300000.0 | NaN | 1.6 |
| 6309 | (500) days of summer | | NaN | 2009 | 2009-07-17 | 7500000.0 | 34439060.0 | 32425665.0 | 4.891301e+07 | 4.6 |
| 1830 | 1,000 times good night | | NaN | 2014 | NaT | NaN | NaN | 53900.0 | NaN | 6.9 |
| 2395 | 10 cloverfield lane | | 38100000.0 | 2016 | 2016-03-11 | 5000000.0 | 108286422.0 | 72091499.5 | 1.374755e+08 | 9.1 |
| 5913 | 10 days in a madhouse | | NaN | 2015 | 2015-11-11 | 12000000.0 | 14616.0 | 14616.0 | 1.878989e+04 | 1.8 |

In [34]:

```
1 # Use plain formatting to remove scientific notation
2 money_df['cpi_worldwide_gross'] = pd.to_numeric(money_df['cpi_worldwide_gross'], errors='coerce')
3 money_df['cpi_production_budget'] = pd.to_numeric(money_df['cpi_production_budget'], errors='coerce')
4 money_df['cpi_foreign_gross'] = pd.to_numeric(money_df['cpi_foreign_gross'], errors='coerce')
5 money_df['cpi_avg Domestic_gross'] = pd.to_numeric(money_df['cpi_avg Domestic_gross'], errors='coerce')
6
7 # Create domestic_profit column
8 money_df['domestic_profit'] = money_df.apply(lambda x: (x['cpi_avg Domestic_gross']) - (x['cpi_production_budget']))
9
10 # Get summary statistics for profit
11 money_df['domestic_profit'].describe().apply(lambda x: format(x, 'f'))
```

Out[34]:

```
count      5238.000000
mean     31169023.887285
std      144654769.502849
min     -366177833.700231
25%    -14361218.380512
50%     397569.143800
75%     39215148.715017
max     4269784084.168346
Name: domestic_profit, dtype: object
```

Profit and loss amounts vary greatly ranging from approximately (in US Dollars) 360 million in losses to 4 billion in profits. The median indicates a movie profit of approximately 400 thousand.

In [35]:

```

1 # Calculate the domestic profit margin for each movie and make a new column
2 money_df['domestic_profit_margin'] = (money_df['domestic_profit'] / money_df['cpi_avg_domestic_gross'])
3
4 print('Total number of results:', len(money_df))
5 display(money_df.head())

```

Total number of results: 7333

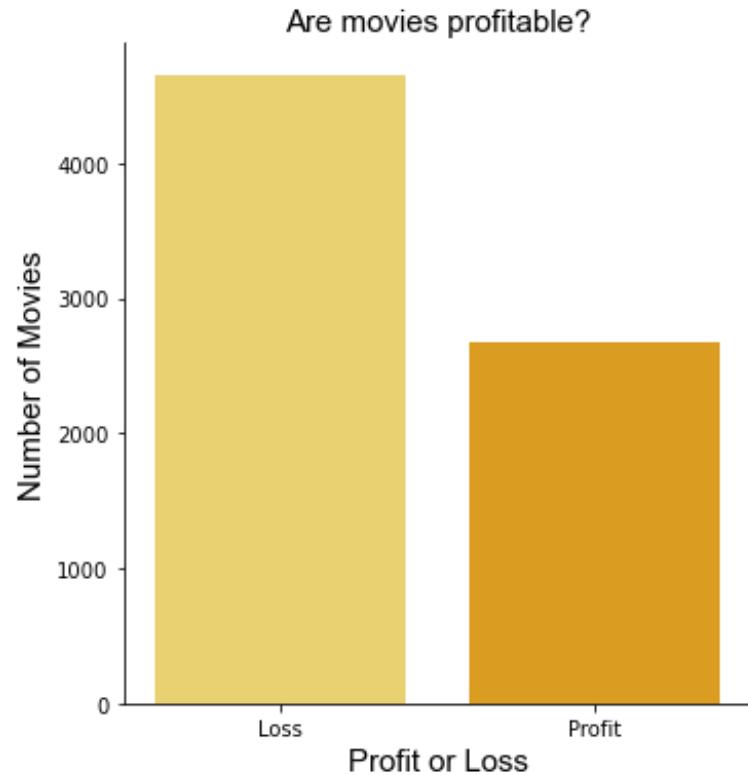
| | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domestic_gross |
|------|------------------------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|------------------------|
| 2093 | '71 | 3550000.0 | 2015 | NaT | NaN | NaN | 1300000.0 | NaN | 1.6 |
| 6309 | (500) days of summer | NaN | 2009 | 2009-07-17 | 7500000.0 | 34439060.0 | 32425665.0 | 4.891301e+07 | 4.6 |
| 1830 | 1,000 times good night | NaN | 2014 | NaT | NaN | NaN | 53900.0 | NaN | 6.9 |
| 2395 | 10 cloverfield lane | 38100000.0 | 2016 | 2016-03-11 | 5000000.0 | 108286422.0 | 72091499.5 | 1.374755e+08 | 9.1 |
| 5913 | 10 days in a madhouse | NaN | 2015 | 2015-11-11 | 12000000.0 | 14616.0 | 14616.0 | 1.878989e+04 | 1.8 |

→ Some of the percentages don't add to 100% so the difference may be attributed to documenting movie gross profits at different times during the movie's runtime. If the value is negative, it may also mean a profit loss.

The number of profitable movies can be visualized using `domestic_profit`=profit margin, as defined above.

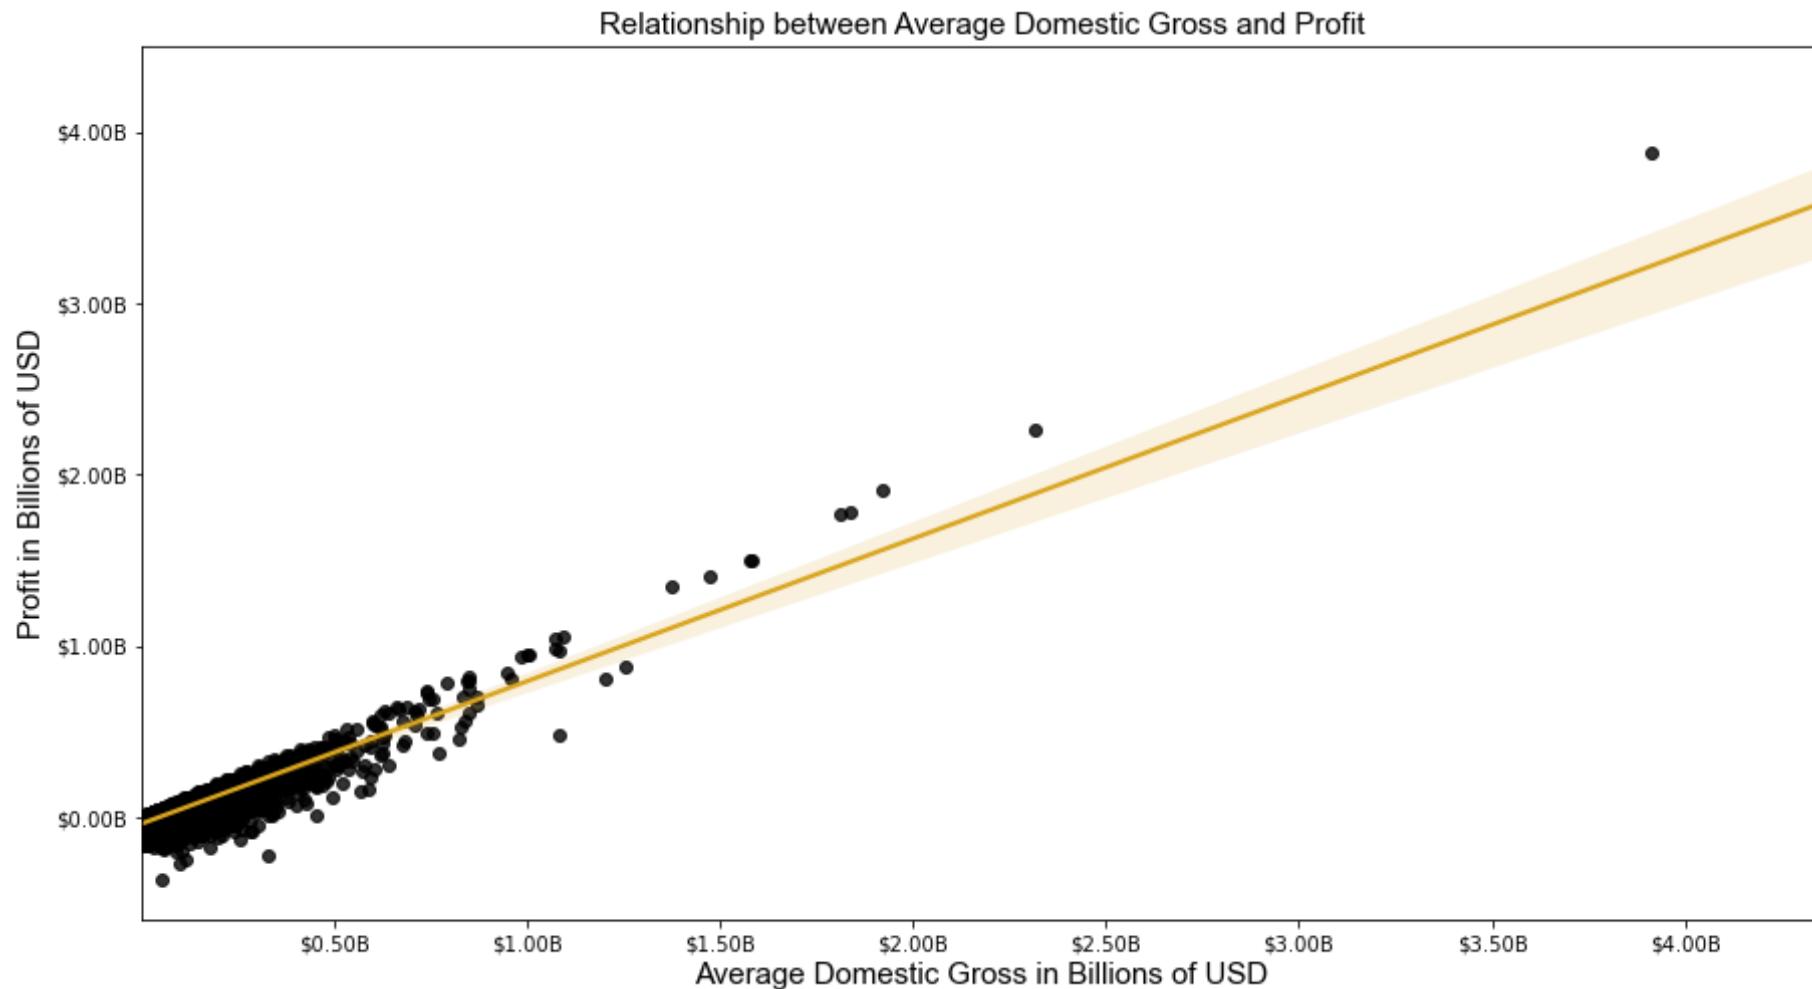
In [36]:

```
1 # Domestic_profit column has both positive and negative values
2 money_df['Profit or Loss'] = money_df['domestic_profit'].apply(lambda x: 'Profit' if x >= 0 else 'Loss')
3
4 # Create plot
5 sns.catplot(x='Profit or Loss', kind='count', palette=['#FDCC5C', '#F9A602'], data=money_df)
6 plt.xlabel('Profit or Loss', fontfamily = 'Arial', fontsize = 15)
7 plt.ylabel('Number of Movies', fontfamily = 'Arial', fontsize = 15)
8 plt.title('Are movies profitable?', fontfamily = 'Arial', fontsize = 15)
9 plt.show()
```



In [37]:

```
1 # Plot relationship between profit and average domestic gross with confidence interval line
2 fig, ax = plt.subplots(figsize=(15, 8))
3
4 # Create plot
5 sns.regplot(
6     x='cpi_avg_domestic_gross',
7     y='domestic_profit',
8     data=money_df,
9     scatter_kws={"color": 'black'},
10    line_kws={"color": '#DAA520'},
11    ci=95, # Adjust the confidence interval
12 )
13
14 plt.title('Relationship between Average Domestic Gross and Profit', fontfamily='Arial', fontsize=15)
15 plt.xlabel('Average Domestic Gross in Billions of USD', fontfamily='Arial', fontsize=15)
16 plt.ylabel('Profit in Billions of USD', fontfamily='Arial', fontsize=15)
17
18 # Define a custom formatter function to display values in billions without scientific notation
19 def billions_formatter(x, pos):
20     """
21     Custom formatter for formatting values in billions without scientific notation.
22
23     Parameters:
24     - x : Value to be formatted
25     - pos : Position of the tick on the axis
26
27     Returns:
28     - formatted_string : Formatted string representing the value in billions
29     """
30     return f'{x / 1e9:.2f}B'
31
32 # Apply the custom formatter to the x-axis and y-axis
33 ax.xaxis.set_major_formatter(billions_formatter)
34 ax.yaxis.set_major_formatter(billions_formatter)
35
36 plt.show()
37
38
```



```
In [38]: 1 # Initial setup of top_profit as top 100 domestic profit movies
2 top_profit = money_df.sort_values('domestic_profit', ascending = False)[:100]
```

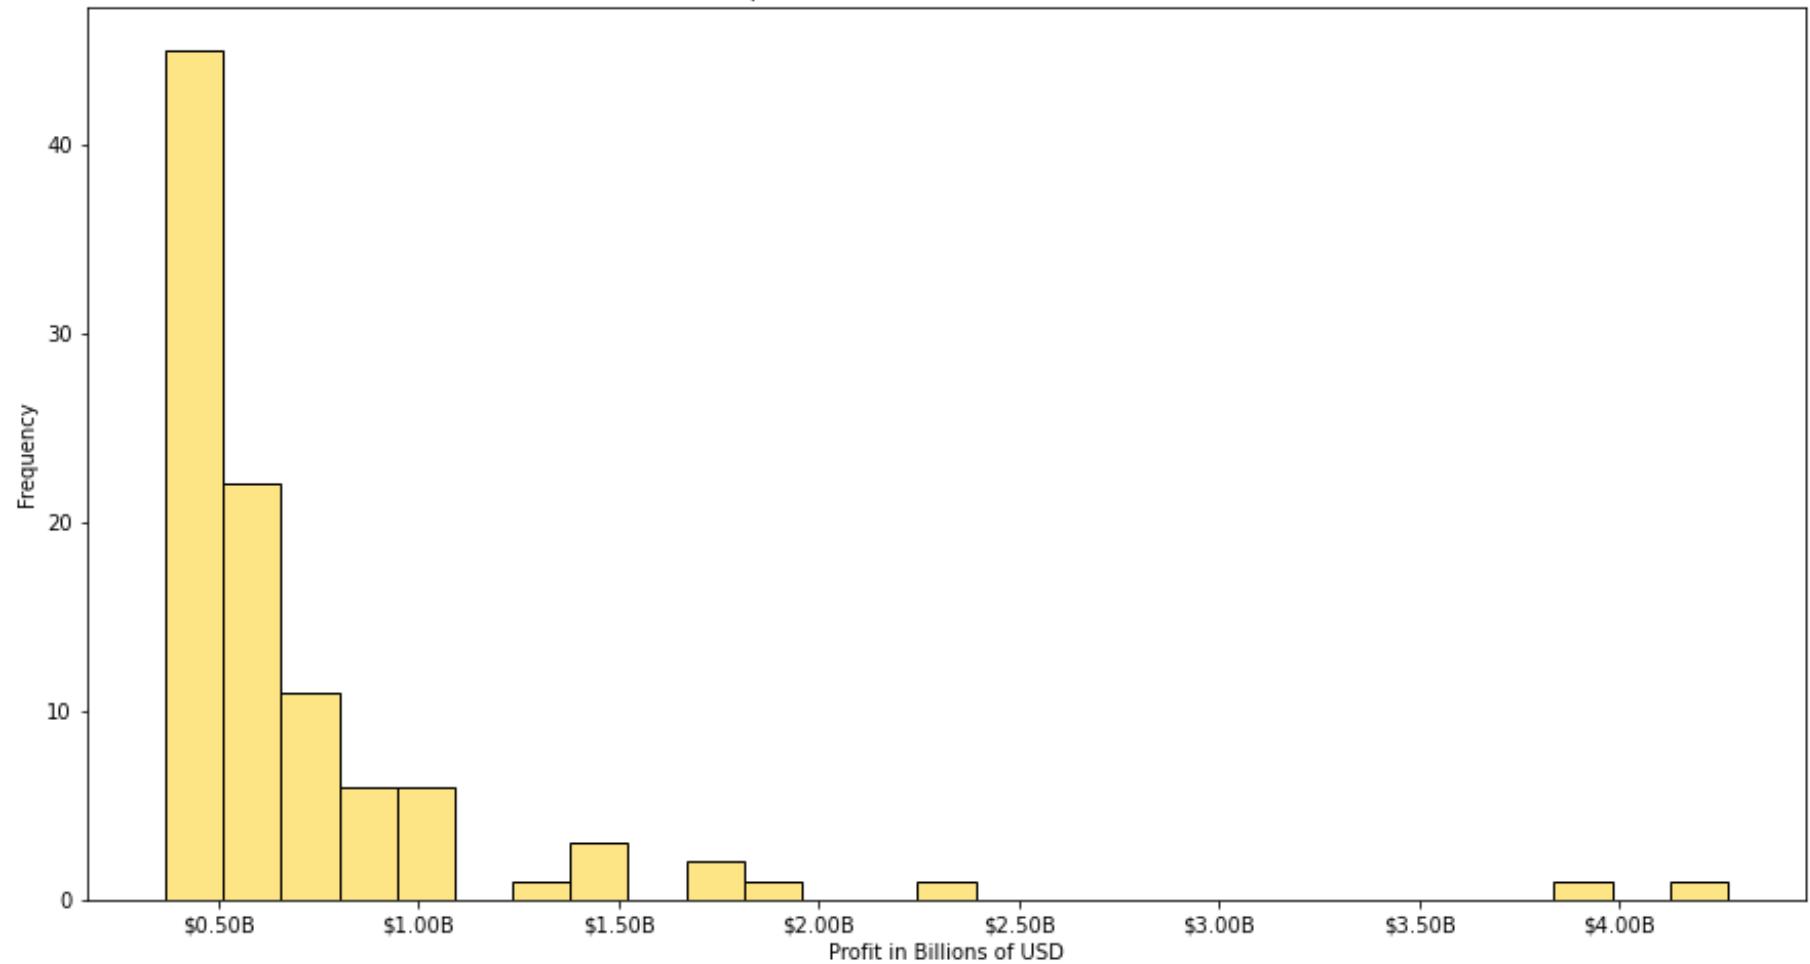
```
In [39]: 1 # Describe profit
2 # Use plain format to remove scientific notation - improves readability
3 top_profit['domestic_profit'].describe().apply(lambda x: format(x, 'f'))
```

```
Out[39]: count      100.000000
mean      721823114.735430
std       602538100.653499
min       367434764.705882
25%      413208392.839609
50%      534583077.617109
75%      758235838.292039
max      4269784084.168346
Name: domestic_profit, dtype: object
```

In [40]:

```
1 # Plot distribution of profit using histplot
2 fig, ax = plt.subplots(figsize=(15, 8))
3
4 sns.histplot(top_profit['domestic_profit'], color= '#FDDC5C', fill=True, edgecolor='black')
5 plt.title('Top 100 Movies Distribution of Profit')
6 plt.xlabel('Profit in Billions of USD')
7 plt.ylabel('Frequency')
8
9 # Define a custom formatter function to display values in billions without scientific notation
10 def billions_formatter(x, pos):
11     """
12         Custom formatter for formatting values in billions without scientific notation.
13
14     Parameters:
15     - x : Value to be formatted
16     - pos : Position of the tick on the axis
17
18     Returns:
19     - formatted_string : Formatted string representing the value in billions
20     """
21     return f'${x / 1e9:.2f}B'
22
23 # Apply the custom formatter to the x-axis and y-axis
24 ax.xaxis.set_major_formatter(billions_formatter)
25
26 plt.show()
```

Top 100 Movies Distribution of Profit



This graph shows that the median profit for top 100 domestic movies is approximately 500 million US dollars.

```
In [41]: 1 # Get summary statistics for profit margin  
2 # Use plain formatting to remove scientific notation  
3 money_df['domestic_profit_margin'].describe().apply(lambda x: format(x, 'f'))
```

```
Out[41]: count      5238.000000  
mean      -4459.698385  
std       70484.653187  
min      -3267873.856209  
25%      -156.738678  
50%       2.931248  
75%      55.418350  
max      99.944444  
Name: domestic_profit_margin, dtype: object
```

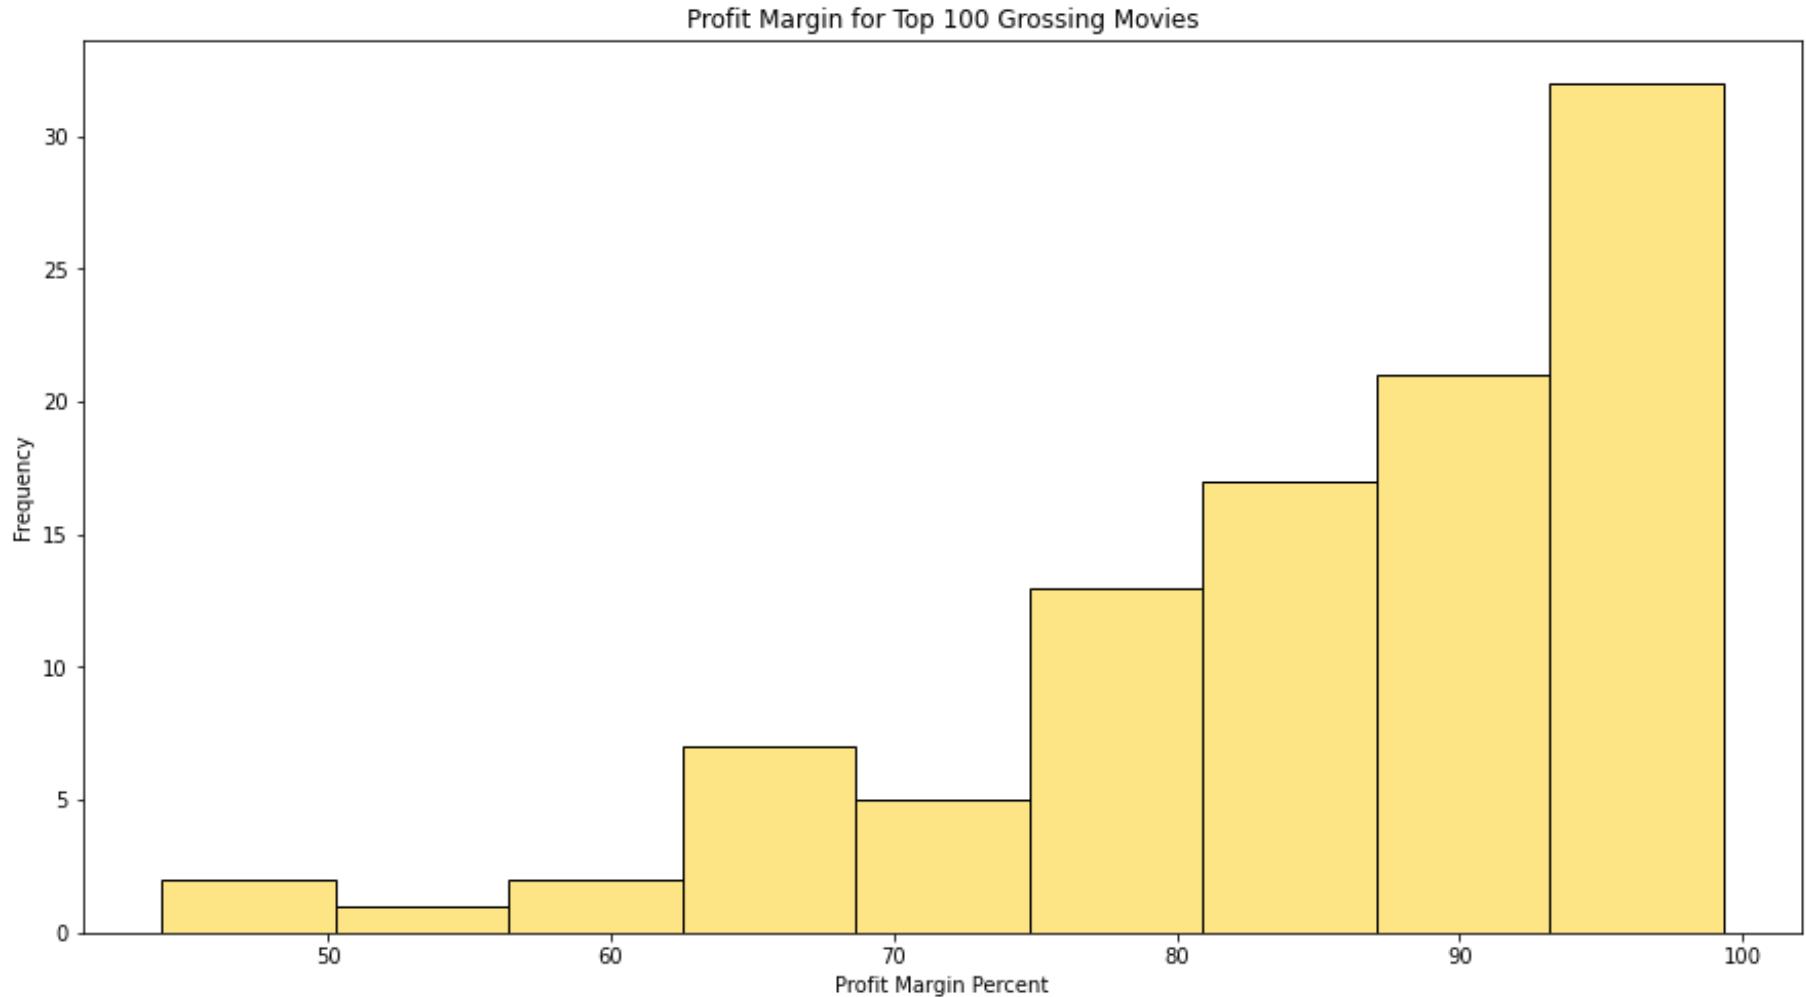
→ There is a significant outlier with a large negative value. Additionally, a considerable portion of these movies have a negative profit margin, indicating a loss. So again, the analysis will only consider the top 100 profitable movies.

```
In [42]: 1 # Set top_profit as top 100 movies  
2 # Describe profit  
3 top_profit['domestic_profit_margin'].describe().apply(lambda x: format(x, 'f'))
```

```
Out[42]: count      100.000000  
mean      85.139087  
std       11.931149  
min      44.116274  
25%      79.593959  
50%      88.051873  
75%      94.760766  
max      99.324348  
Name: domestic_profit_margin, dtype: object
```

In [43]:

```
1 # Plot distribution of profit margin for top_profit using histplot
2 plt.figure(figsize = (15,8))
3 sns.histplot(top_profit['domestic_profit_margin'], color= '#FDDC5C', fill=True, edgecolor='black')
4 plt.title('Profit Margin for Top 100 Grossing Movies')
5 plt.xlabel('Profit Margin Percent')
6 plt.ylabel('Frequency')
7 plt.show()
8
```

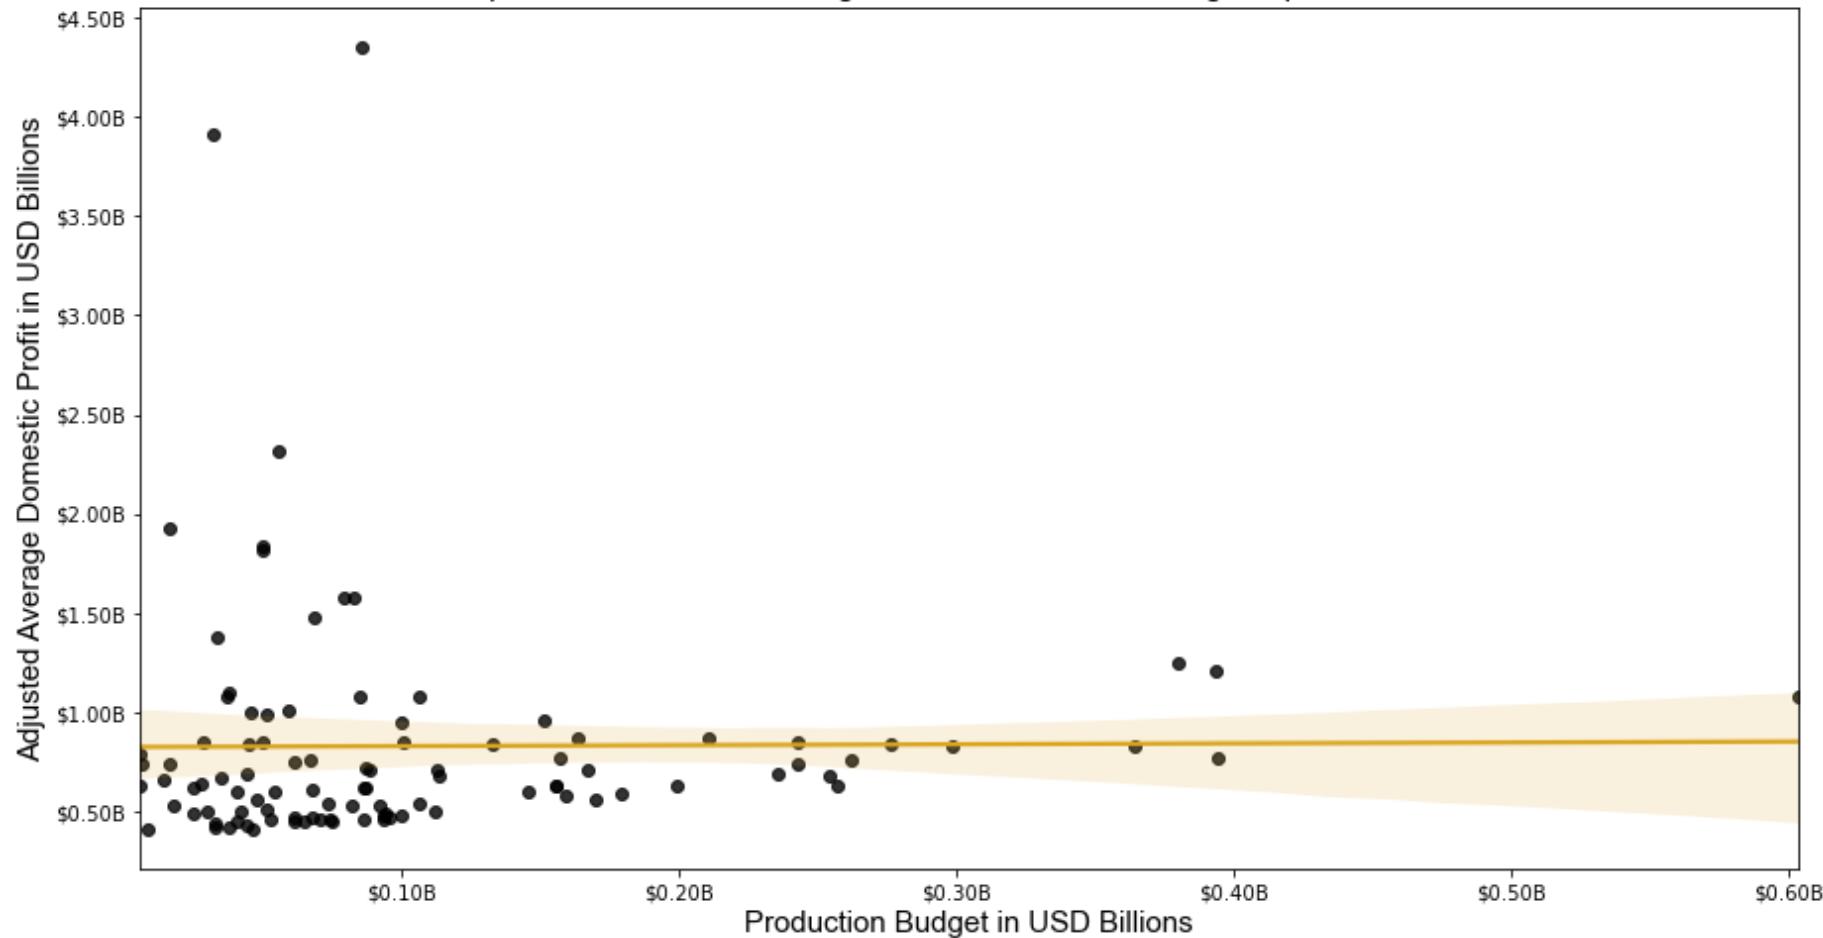


The median profit margin for a top 100 all-time domestic grossing movies is approximately 88%

In [44]:

```
1 # Create a scatter plot with confidence interval line
2 fig, ax = plt.subplots(figsize=(15, 8))
3
4 sns.regplot(
5     x='cpi_production_budget',
6     y='cpi_avg_domestic_gross',
7     data=top_profit,
8     scatter_kws={"color": 'black'},
9     line_kws={"color": '#DAA520'},
10    ci=95, # Adjust the confidence interval
11 )
12
13 plt.title('Relationship between Production Budget and Domestic Profit Margin Top 100 Profitable Movies')
14 plt.xlabel('Production Budget in USD Billions', fontfamily='Arial', fontsize=15)
15 plt.ylabel('Adjusted Average Domestic Profit in USD Billions', fontfamily='Arial', fontsize=15)
16
17 # Define a custom formatter function to display values in billions without scientific notation
18 def billions_formatter(x, pos):
19     """
20         Custom formatter for formatting values in billions without scientific notation.
21
22     Parameters:
23     - x : Value to be formatted
24     - pos : Position of the tick on the axis
25
26     Returns:
27     - formatted_string : Formatted string representing the value in billions
28     """
29     return f'{x / 1e9:.2f}B'
30
31 # Apply the custom formatter to both x-axis and y-axis
32 x_formatter = FuncFormatter(billions_formatter)
33 ax.xaxis.set_major_formatter(x_formatter)
34 ax.yaxis.set_major_formatter(billions_formatter)
35
36 plt.show()
37
```

Relationship between Production Budget and Domestic Profit Margin Top 100 Profitable Movies



```
In [45]: 1 # Set top_profit as top 100 movies
  2 # Describe profit
  3 top_profit['cpi_production_budget'].describe().apply(lambda x: format(x, 'f'))
```

```
Out[45]: count      100.000000
mean    106905490.695924
std     101674087.339997
min     5332285.000000
25%    44218569.249904
50%    74142873.566314
75%    136047153.909770
max     603617790.870572
Name: cpi_production_budget, dtype: object
```

This plot between production budget and adjusted average domestic gross of the top 100 grossing movies, does not show a clear correlation. It would seem that movies have seen success with small budgets and others despite robust budgets, failed to perform. It appears that budgets between 44 and 136 million seem to be reasonable.

In [46]:

```
1 # Get top profit movies in descending order
2 print('Total number of results:', len(top_profit))
3 display(top_profit.head())
```

Total number of results: 100

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_dome |
|------|--|-------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|--------------|
| 6814 | gone with the wind | | NaN | 1939 | 1939-12-15 | 3900000.0 | 390525192.0 | 198680470.0 | 8.560706e+09 | 4.3% |
| 7293 | snow white and the seven dwarfs | | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.9% |
| 5928 | star wars ep. iv: a new hope | | NaN | 1977 | 1977-05-25 | 11000000.0 | 786598007.0 | 460998007.0 | 3.955082e+09 | 2.3% |
| 7497 | bambi | | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.9% |
| 7073 | pinocchio | | NaN | 1940 | 1940-02-09 | 2289247.0 | 84300000.0 | 84300000.0 | 1.834741e+09 | 1.8% |

In [47]:

```
1 # Get top profit margin movies
2 top_margin = top_profit.sort_values('domestic_profit_margin', ascending = False)
3 print('Total number of results:', len(top_margin))
4 display(top_margin.head())
```

Total number of results: 100

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domes |
|------|---|-------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|---------------|
| 7521 | american graffiti | | NaN | 1973 | 1973-08-11 | 777000.0 | 140000000.0 | 115000000.0 | 9.607721e+08 | 7.89 |
| 7293 | snow white and the seven dwarfs | | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.91 |
| 7507 | billy jack | | NaN | 1971 | 1971-01-01 | 800000.0 | 98000000.0 | 98000000.0 | 7.373036e+08 | 7.37 |
| 7497 | bambi | | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.92 |
| 7384 | rocky | | NaN | 1976 | 1976-11-21 | 1000000.0 | 225000000.0 | 117235147.0 | 1.204885e+09 | 6.27 |

The movies with the all-time highest profit differ from those with the highest all-time profit margin.

Profit and Profit Margin Analysis Summary

Production budgets for the top 100 profit movies varies greatly ranging from around 5 million to around 600 million. This shows that increasing the movie production budget does not necessarily increase profits.

A further investigation of the aspects included in the production budget could reveal how to best allocate the production budget.

Aggregating additional release date and title names

In [48]: 1 tmdb_movies.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        26517 non-null   int64  
 1   genre_ids         26517 non-null   object  
 2   id                26517 non-null   int64  
 3   original_language 26517 non-null   object  
 4   original_title    26517 non-null   object  
 5   popularity        26517 non-null   float64 
 6   release_date      26517 non-null   object  
 7   title              26517 non-null   object  
 8   vote_average      26517 non-null   float64 
 9   vote_count         26517 non-null   int64  
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

In [49]:

```

1 # Change the movie titles to lowercase
2 tmdb_movies['original_title'] = tmdb_movies['original_title'].str.lower()
3 tmdb_movies['title'] = tmdb_movies['title'].str.lower()
4
5 # Create year column
6 tmdb_movies['year'] = pd.DatetimeIndex(tmdb_movies['release_date']).year
7
8 print('Total number of results:', len(tmdb_movies))
9 display(tmdb_movies.head())

```

Total number of results: 26517

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | vote_average | vote_count | year |
|---|---------------|------------------------|-------|-------------------|--|------------|--------------|--|--------------|------------|------|
| 0 | 0 | [12, 14, 10751] | 12444 | en | harry potter and the deathly hallows: part 1 | 33.533 | 2010-11-19 | harry potter and the deathly hallows: part 1 | 7.7 | 10788 | 2010 |
| 1 | 1 | [14, 12, 16, 10751] | 10191 | en | how to train your dragon | 28.734 | 2010-03-26 | how to train your dragon | 7.7 | 7610 | 2010 |
| 2 | 2 | [12, 28, 878] | 10138 | en | iron man 2 | 28.515 | 2010-05-07 | iron man 2 | 6.8 | 12368 | 2010 |
| 3 | 3 | [16, 35, 10751] | 862 | en | toy story | 28.005 | 1995-11-22 | toy story | 7.9 | 10174 | 1995 |
| 4 | 4 | [28, 878, 12] | 27205 | en | inception | 27.920 | 2010-07-16 | inception | 8.3 | 22186 | 2010 |

In [50]:

```
1 # Merge money_df and tmdb_movies dataframes
2 many_titles_df = pd.merge(money_df, tmdb_movies, on=['title', 'year'], how='outer')
3 print('Total number of results:', len(many_titles_df))
4 display(many_titles_df.head())
```

Total number of results: 31373

| | | title | foreign_gross | year | release_date_x | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domes |
|---|---------------------------------|-------|---------------|------|----------------|-------------------|-----------------|--------------------|---------------------|---------------|
| 0 | '71 | | 3550000.0 | 2015 | NaT | NaN | NaN | 1300000.0 | NaN | 1.67 |
| 1 | (500) days of summer | | NaN | 2009 | 2009-07-17 | 7500000.0 | 34439060.0 | 32425665.0 | 4.891301e+07 | 4.60 |
| 2 | 1,000 times good night | | NaN | 2014 | NaT | NaN | NaN | 53900.0 | NaN | 6.93 |
| 3 | 10 cloverfield lane | | 38100000.0 | 2016 | 2016-03-11 | 5000000.0 | 108286422.0 | 72091499.5 | 1.374755e+08 | 9.15 |

In [51]:

```

1 # Fill NaN or 0.0 cells in column release_date_y with value from release_date_x
2 many_titles_df['release_date_y'].fillna(many_titles_df['release_date_x'], inplace=True)
3
4 # Drop the 'release_date_x' column
5 many_titles_df.drop(columns=['release_date_x'], inplace=True)
6
7 # Rename 'release_date_y' to 'release_date'
8 many_titles_df.rename(columns={'release_date_y': 'release_date'}, inplace=True)
9
10 # Drop the 'id' and 'Unnamed columns' columns
11 many_titles_df.drop(columns=['id', 'Unnamed: 0'], inplace=True)
12
13 # Drop duplicates
14 many_titles_df.drop_duplicates()
15
16 print('Total number of results:', len(many_titles_df))
17 display(many_titles_df.head())

```

Total number of results: 31373

| | | title | foreign_gross | year | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domestic_gross | cpi_fc |
|---|---------------------------------|-------|---------------|------|-------------------|-----------------|--------------------|---------------------|------------------------|--------------|
| 0 | '71 | | 3550000.0 | 2015 | | NaN | NaN | 1300000.0 | NaN | 1.671241e+06 |
| 1 | (500) days of summer | | | NaN | 2009 | 7500000.0 | 34439060.0 | 32425665.0 | 4.891301e+07 | 4.605343e+07 |
| 2 | 1,000 times good night | | | NaN | 2014 | NaN | NaN | 53900.0 | NaN | 6.937448e+04 |
| 3 | 10 cloverfield lane | | 38100000.0 | 2016 | | 5000000.0 | 108286422.0 | 72091499.5 | 1.374755e+08 | 9.152410e+07 |
| 4 | 10 days in a madhouse | | | NaN | 2015 | 12000000.0 | 14616.0 | 14616.0 | 1.878989e+04 | 1.878989e+04 |

```
In [52]: 1 # Replace NaN values in 'original_title' with values from 'title'  
2 many_titles_df['original_title'].fillna(many_titles_df['title'], inplace=True)  
3  
4 # Replace NaN values in 'title' with values from 'original_title'  
5 many_titles_df['title'].fillna(many_titles_df['original_title'], inplace=True)  
6
```

```
In [53]: 1 # Find and count duplicates between the two columns  
2 duplicate_count = many_titles_df[many_titles_df.duplicated(subset=['original_title', 'title'], keep=False)]  
3  
4 print(f"Number of duplicate titles between 'original_title' and 'title': {duplicate_count}")  
5
```

Number of duplicate titles between 'original_title' and 'title': 3924

In [54]:

```
1 # Create a boolean mask to identify rows where values in the two columns differ
2 mask = many_titles_df['original_title'] != many_titles_df['title']
3
4 # Extract rows with differences in the specified columns
5 differing_rows = many_titles_df[mask][['original_title', 'title']]
6
7 # Print the differing rows
8 print("Rows with differences in 'original_title' and 'title':")
9 display(differing_rows)
10
11 # Print the length of the list
12 f"Number of differing rows: {len(differing_rows)}"
```

Rows with differences in 'original_title' and 'title':

| | original_title | title |
|-------|------------------------------|---------------------------------|
| 18 | 十三人の刺客 | 13 assassins |
| 21 | elser | 13 minutes |
| 29 | 辛亥革命 | 1911 |
| 36 | हू स्टेट्स | 2 states |
| 103 | 7 cajas | 7 boxes |
| ... | ... | ... |
| 31277 | czarne lusterko: 69.90 | little black mirror: 69.90 |
| 31292 | el verano del león eléctrico | the summer of the electric lion |
| 31301 | contes de juillet | july tales |
| 31353 | la última virgen | the last virgin |
| 31365 | dreamaway | dream away |

2532 rows × 2 columns

Out[54]: 'Number of differing rows: 2532'

```
In [55]: 1 # Drop the 'original_title' column  
2 many_titles_df.drop(columns=['original_title'], inplace=True)
```

→ Code blocks 54-56 discover that the differences between the title names can be attributed to being documented in different languages.

Extract and aggregate SQLite Tables to Pandas dataframe

```
In [56]: 1 # Connect to SQL Database  
2 conn = sqlite3.connect('Data/im.db')  
3  
4 # Open SQL Table  
5 table_names = ['movie_basics']  
6  
7 # Create an empty dictionary to store DataFrames  
8 SQL_dataframes = {}  
9  
10 # Loop through table names and read data into DataFrames  
11 for table in table_names:  
12     sql_query = f'SELECT * FROM {table}'  
13     SQL_dataframes[table] = pd.read_sql(sql_query, conn)  
14
```

In [57]:

```
1 # Preview the data from SQL DataFrames
2 for name, df in SQL_dataframes.items():
3     print(name)
4     print('Total number of results:', len(df))
5     display(df.head()) # Using display instead of print leads to neater formatting in Jupyter Notebook
6     print(df.info())
7     print('----- \n')
8
```

movie_basics

Total number of results: 146144

| | movie_id | primary_title | original_title | start_year | runtime_minutes | genres |
|---|-----------|---------------------------------|----------------------------|------------|-----------------|----------------------|
| 0 | tt0063540 | Sunghursh | Sunghursh | 2013 | 175.0 | Action,Crime,Drama |
| 1 | tt0066787 | One Day Before the Rainy Season | Ashad Ka Ek Din | 2019 | 114.0 | Biography,Drama |
| 2 | tt0069049 | The Other Side of the Wind | The Other Side of the Wind | 2018 | 122.0 | Drama |
| 3 | tt0069204 | Sabse Bada Sukh | Sabse Bada Sukh | 2018 | NaN | Comedy,Drama |
| 4 | tt0100275 | The Wandering Soap Opera | La Telenovela Errante | 2017 | 80.0 | Comedy,Drama,Fantasy |

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie_id        146144 non-null   object 
 1   primary_title   146144 non-null   object 
 2   original_title  146123 non-null   object 
 3   start_year      146144 non-null   int64  
 4   runtime_minutes 114405 non-null   float64
 5   genres          140736 non-null   object 
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
None
```

In [58]:

```

1 # Access the 'movie_basics' DataFrame from the SQL_dataframes dictionary
2 movie_basics_df = SQL_dataframes['movie_basics']
3
4 # Create a copy of the DataFrame to avoid altering the original
5 movie_basics_ga = movie_basics_df.copy()
6
7 # Merge the DataFrames on 'primary_title' and 'start_year' in 'movie_basics_ga' with 'title' and 'year'
8 main_df = pd.merge(
9     movie_basics_ga,
10    many_titles_df,
11    left_on=['primary_title', 'start_year'],
12    right_on=['title', 'year'],
13    how='outer'
14 ).drop_duplicates()
15
16
17
18 print('Total number of results:', len(main_df))
19 display(main_df.head())

```

Total number of results: 176449

| | movie_id | primary_title | original_title | start_year | runtime_minutes | genres | title | foreign_gross | year | production_budget | ... | cpi |
|---|-----------|---------------------------------|----------------------------|------------|-----------------|----------------------|-------|---------------|------|-------------------|-----|-----|
| 0 | tt0063540 | Sunghursh | Sunghursh | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN | ... | NaN |
| 1 | tt0066787 | One Day Before the Rainy Season | Ashad Ka Ek Din | 2019.0 | 114.0 | Biography,Drama | NaN | NaN | NaN | NaN | ... | NaN |
| 2 | tt0069049 | The Other Side of the Wind | The Other Side of the Wind | 2018.0 | 122.0 | Drama | NaN | NaN | NaN | NaN | ... | NaN |
| 3 | tt0069204 | Sabse Bada Sukh | Sabse Bada Sukh | 2018.0 | NaN | Comedy,Drama | NaN | NaN | NaN | NaN | ... | NaN |
| 4 | tt0100275 | The Wandering Soap Opera | La Telenovela Errante | 2017.0 | 80.0 | Comedy,Drama,Fantasy | NaN | NaN | NaN | NaN | ... | NaN |

5 rows × 25 columns

In [59]:

```
1 # Convert strings to lowercase in 'primary_title_y' and 'title' columns
2 main_df['primary_title'] = main_df['primary_title'].str.lower()
3
4 # Convert strings to lowercase in 'original_title' and 'title' columns
5 main_df['original_title'] = main_df['original_title'].str.lower()
6
7 # Find and count duplicates between the two columns
8 duplicate_count = main_df[main_df.duplicated(subset=['primary_title', 'title'], keep=False)].shape[0]
9
10 print(f"Number of duplicate titles between 'primary_title' and 'title': {duplicate_count}")
```

Number of duplicate titles between 'primary_title' and 'title': 19165

In [60]:

```
1 # Fill NaN or 0.0 cells in column primary_title with value from title
2 main_df['primary_title'].fillna(main_df['title'], inplace=True)
3
4 # Fill NaN or 0.0 cells in column primary_title with value from original_title
5 main_df['primary_title'].fillna(main_df['original_title'], inplace=True)
6
7 # Fill NaN or 0.0 cells in column start_year with value from year
8 main_df['start_year'].fillna(main_df['year'], inplace=True)
9
10 # Drop the 'year' column
11 main_df.drop(columns=['year'], inplace=True)
12
13 # Rename 'start_year' to 'years'
14 main_df.rename(columns={'start_year': 'release_year'}, inplace=True)
15
16 # Drop the 'title' column
17 main_df.drop(columns=['title', 'original_title'], inplace=True)
18
19 # Rename 'primary_title' to 'title'
20 main_df.rename(columns={'primary_title': 'title'}, inplace=True)
21
22 print('Total number of results:', len(main_df))
23 display(main_df.head())
```

Total number of results: 176449

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_ |
|---|-----------|--|--------------|-----------------|----------------------|---------------|-------------------|-----------------|---------------|
| 0 | tt0063540 | sunghursh one day before the rainy season | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN |
| 1 | tt0066787 | the other side of the wind | 2019.0 | 114.0 | Biography,Drama | NaN | NaN | NaN | NaN |
| 2 | tt0069049 | sabse bada sukh | 2018.0 | 122.0 | Drama | NaN | NaN | NaN | NaN |
| 3 | tt0069204 | the wandering soap opera | 2018.0 | NaN | Comedy,Drama | NaN | NaN | NaN | NaN |
| 4 | tt0100275 | | 2017.0 | 80.0 | Comedy,Drama,Fantasy | NaN | NaN | NaN | NaN |

5 rows × 22 columns

In [61]:

```
1 # Check for duplicates in the 'title' column
2 title_year_duplicates = main_df.duplicated(subset=['title','release_year'], keep=False)
3
4 # Display a sample of rows with duplicate titles (view only)
5 sample_duplicates = main_df[title_year_duplicates].sort_values(by='title')
6
7 print('Total number of results:', len(sample_duplicates))
8 display(sample_duplicates.head())
```

Total number of results: 29180

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_g |
|--------|-----------|---------------|--------------|-----------------|-------------|---------------|-------------------|-----------------|----------------|
| 168121 | NaN | #allmymovies | 2015.0 | NaN | NaN | NaN | NaN | NaN | NaN |
| 128223 | tt7853996 | #allmymovies | 2015.0 | NaN | Documentary | NaN | NaN | NaN | NaN |
| 173604 | NaN | #captured | 2017.0 | NaN | NaN | NaN | NaN | NaN | NaN |
| 115921 | tt6856592 | #captured | 2017.0 | 81.0 | Thriller | NaN | NaN | NaN | NaN |
| 170488 | NaN | #followfriday | 2016.0 | NaN | NaN | NaN | NaN | NaN | NaN |

5 rows × 22 columns

In [62]:

```

1 # Split out main_df into duplicates and unique
2 # Make 'unique' title df
3 main_df_unique = main_df[~title_year_duplicates].copy()
4
5 # Handle duplicates and concat them back to the above unique df
6 # Make 'duplicate' title df
7 main_df_dups = main_df[title_year_duplicates].copy()
8
9 print('Total number of results:', len(main_df_dups))
10 display(main_df_dups.head())

```

Total number of results: 29180

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_gross |
|----|-----------|----------------------------|--------------|-----------------|------------------------|---------------|-------------------|-----------------|--------------------|
| 2 | tt0069049 | the other side of the wind | 2018.0 | 122.0 | Drama | NaN | NaN | NaN | NaN |
| 17 | tt0192528 | heaven & hell | 2018.0 | 104.0 | Drama | NaN | NaN | NaN | NaN |
| 21 | tt0250404 | godfather | 2012.0 | NaN | Crime,Drama | NaN | NaN | NaN | NaN |
| 22 | tt3270104 | godfather | 2012.0 | 145.0 | Drama | NaN | NaN | NaN | NaN |
| 24 | tt0255820 | return to babylon | 2013.0 | 75.0 | Biography,Comedy,Drama | NaN | NaN | NaN | NaN |

5 rows × 22 columns

In [63]:

```

1 # Get value counts
2 title_counts = main_df_dups['title'].value_counts()
3
4 # Filter titles that appear more than twice
5 titles_more_than_twice = title_counts[title_counts > 2].index.tolist()
6
7 # Display the titles
8 print("Titles listed more than twice:")
9 print(len(titles_more_than_twice))

```

Titles listed more than twice:

805

→ It looks like some of the films may have multiple release years, which is a factor to consider in removing duplicates.

In [64]:

```
1 # Define a custom function to choose the longer of the two genres
2 def choose_longer_genre(series):
3     """
4         All NaN values within genres will be filled from duplicate title entries in the data frame.
5         The longer string value of the two genres is chosen when there are different values between rows.
6     -
7         Input:
8             series : Pandas Series, represents the 'genres' column
9         -
10        Output:
11            value : Chosen genre value based on conditions
12        """
13
14    if series.isnull().any():
15        # If one of the values is missing, choose the non-null value
16        return next((value for value in series.dropna()), None)
17    else:
18        # If both values are present, choose the longer one
19        return max(series, key=len)
20
21 # Apply the custom function to the 'genres' column
22 main_df_dups['genres'] = main_df_dups.groupby(['title', 'release_year'])['genres'].transform(choose_longer_genre)
23
24 # Columns to fill NaN values from the next row
25 cols_to_fill = ['runtime_minutes', 'release_date', 'movie_id', 'cpi_avg_domestic_gross']
26
27 # Forward fill NaN values from the next row
28 main_df_dups[cols_to_fill] = main_df_dups.groupby(['title', 'release_year'])[cols_to_fill].transform(lambda x: x.fillna(x.shift()))
29
30 print('Total number of results:', len(main_df_dups))
31 display(main_df_dups.head())
32
```

Total number of results: 29180

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_gross |
|----|-----------|----------------------------|--------------|-----------------|------------------------|---------------|-------------------|-----------------|--------------------|
| 2 | tt0069049 | the other side of the wind | 2018.0 | 122.0 | Drama | NaN | NaN | NaN | NaN |
| 17 | tt0192528 | heaven & hell | 2018.0 | 104.0 | Drama | NaN | NaN | NaN | NaN |
| 21 | tt0250404 | godfather | 2012.0 | 145.0 | Crime,Drama | NaN | NaN | NaN | NaN |
| 22 | tt3270104 | godfather | 2012.0 | 145.0 | Crime,Drama | NaN | NaN | NaN | NaN |
| 24 | tt0255820 | return to babylon | 2013.0 | 75.0 | Biography,Comedy,Drama | NaN | NaN | NaN | NaN |

5 rows × 22 columns

In [65]:

```

1 # Drop duplicate rows and keep the first occurrence
2 main_df_dups_dropped = main_df_dups.drop_duplicates(subset=['title', 'release_year'], keep='first')
3
4 print('Total number of results:', len(main_df_dups_dropped))
5 display(main_df_dups_dropped.head())

```

Total number of results: 14252

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_gross |
|----|-----------|----------------------------|--------------|-----------------|------------------------|---------------|-------------------|-----------------|--------------------|
| 2 | tt0069049 | the other side of the wind | 2018.0 | 122.0 | Drama | NaN | NaN | NaN | NaN |
| 17 | tt0192528 | heaven & hell | 2018.0 | 104.0 | Drama | NaN | NaN | NaN | NaN |
| 21 | tt0250404 | godfather | 2012.0 | 145.0 | Crime,Drama | NaN | NaN | NaN | NaN |
| 24 | tt0255820 | return to babylon | 2013.0 | 75.0 | Biography,Comedy,Drama | NaN | NaN | NaN | NaN |
| 39 | tt0315642 | wazir | 2016.0 | 103.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN |

5 rows × 22 columns

In [66]:

```

1 # Rebuilding the main_df, but with the duplicates removed
2 # Main_df_unique + main_df_dups_dropped
3 main_df_clean = main_df_unique.append(main_df_dups_dropped, ignore_index=True)
4
5 print('Total number of results:', len(main_df_clean))
6 display(main_df_clean.head())

```

Total number of results: 161521

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_ |
|---|-----------|--|--------------|-----------------|----------------------|---------------|-------------------|-----------------|---------------|
| 0 | tt0063540 | sunghursh | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN |
| 1 | tt0066787 | one day before the rainy season | 2019.0 | 114.0 | Biography,Drama | NaN | NaN | NaN | NaN |
| 2 | tt0069204 | sabse bada sukh | 2018.0 | NaN | Comedy,Drama | NaN | NaN | NaN | NaN |
| 3 | tt0100275 | the wandering soap opera | 2017.0 | 80.0 | Comedy,Drama,Fantasy | NaN | NaN | NaN | NaN |
| 4 | tt0111414 | a thin life | 2018.0 | 75.0 | Comedy | NaN | NaN | NaN | NaN |

5 rows × 22 columns

Genre Analysis

The genre associated with a movie may greatly influence profitability. By analyzing the adjusted average domestic gross for each individual genre, it may shed light on a factor that contributes to a film's success at the box office.

In [67]:

```
1 # Create a copy of the DataFrame to avoid altering the original
2 genre_analysis_explode = main_df_clean.copy()
3
4 # Split the genres into separate rows and create a new 'genre' column
5 genre_analysis_explode['genre'] = genre_analysis_explode['genres'].str.split(',')
6
7 # Explode the 'genre' column to create new rows for each genre
8 genre_analysis_explode = genre_analysis_explode.explode('genre')
9
10 # Display the updated DataFrame
11 print('Total number of results:', len(genre_analysis_explode))
12 display(genre_analysis_explode.head())
13
```

Total number of results: 249658

| | movie_id | title | release_year | runtime_minutes | genres | foreign_gross | production_budget | worldwide_gross | avg_domestic_gross | avg_imdb_rating |
|---|-----------|---------------------------------|--------------|-----------------|--------------------|---------------|-------------------|-----------------|--------------------|-----------------|
| 0 | tt0063540 | sunghursh | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN | NaN |
| 0 | tt0063540 | sunghursh | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN | NaN |
| 0 | tt0063540 | sunghursh | 2013.0 | 175.0 | Action,Crime,Drama | NaN | NaN | NaN | NaN | NaN |
| 1 | tt0066787 | one day before the rainy season | 2019.0 | 114.0 | Biography,Drama | NaN | NaN | NaN | NaN | NaN |
| 1 | tt0066787 | one day before the rainy season | 2019.0 | 114.0 | Biography,Drama | NaN | NaN | NaN | NaN | NaN |

In [68]:

```
1 # Count the occurrences of each genre and convert to DataFrame
2 genre_counts = genre_analysis_explode['genre'].value_counts().reset_index()
3 genre_counts.columns = ['genre', 'count']
4
5 # Get the total sum of value counts
6 total_count = genre_counts['count'].sum()
7
8 # Calculate the percentage of each genre
9 genre_percentages = (genre_counts['count'] / total_count) * 100
10
11 # Create a DataFrame to display the results
12 percent_genre = pd.DataFrame({
13     'Genre': genre_counts['genre'],
14     'Count': genre_counts['count'],
15     'Percentage': genre_percentages
16 })
17
18 # Display the results
19 print('Total number of results:', total_count)
20 display(percent_genre)
21
```

Total number of results: 226988

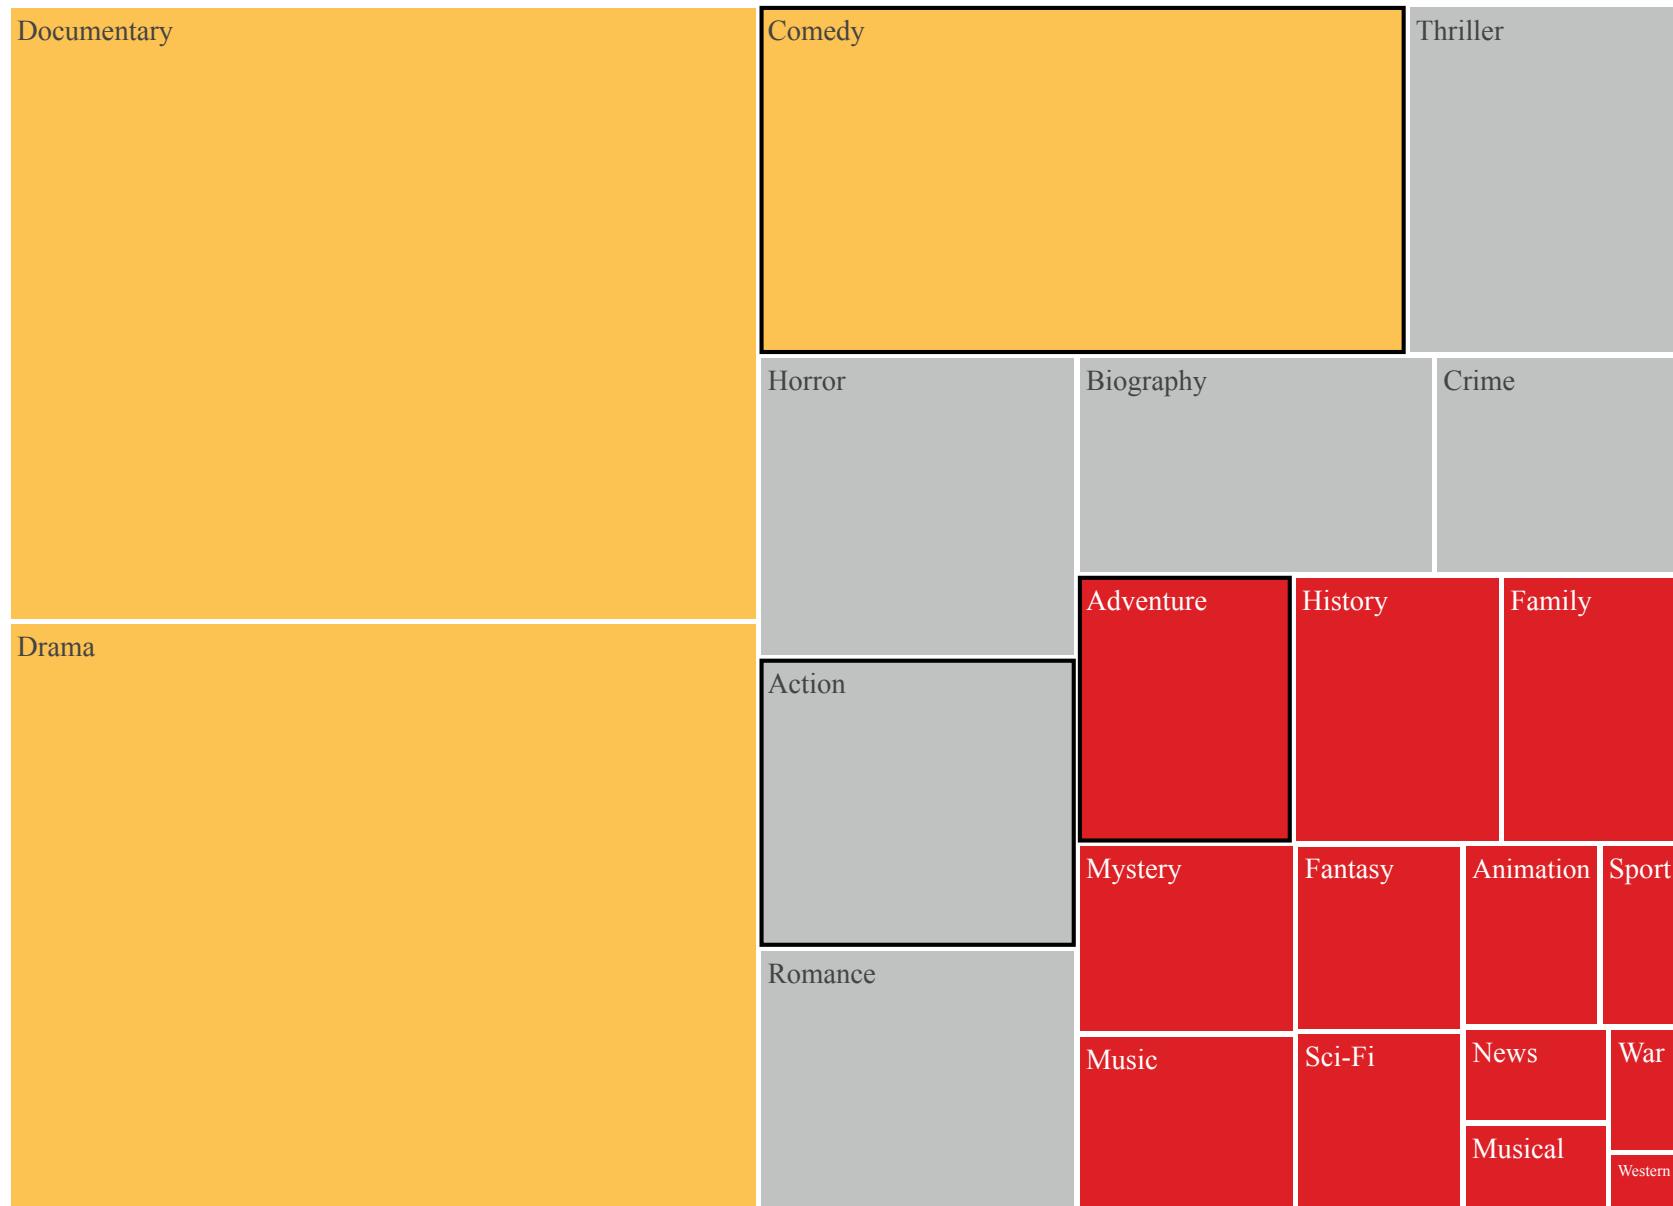
| | Genre | Count | Percentage |
|---|-------------|-------|------------|
| 0 | Documentary | 51191 | 22.552294 |
| 1 | Drama | 48979 | 21.577793 |
| 2 | Comedy | 25120 | 11.066664 |
| 3 | Thriller | 11701 | 5.154898 |
| 4 | Horror | 10655 | 4.694081 |
| 5 | Action | 10222 | 4.503322 |
| 6 | Romance | 9274 | 4.085679 |
| 7 | Biography | 8687 | 3.827075 |
| 8 | Crime | 6667 | 2.937160 |
| 9 | Adventure | 6426 | 2.830987 |

In [69]:

```
1 # Define your color codes for each genre
2 color_codes = {
3     'Documentary': '#FCC252',
4     'Drama': '#FCC252',
5     'Comedy': '#FCC252',
6     'Thriller': '#C0C1C1',
7     'Horror': '#C0C1C1',
8     'Action': '#C0C1C1',
9     'Romance': '#C0C1C1',
10    'Biography': '#C0C1C1',
11    'Crime': '#C0C1C1',
12    'Adventure': '#DC2026',
13    'History': '#DC2026',
14    'Family': '#DC2026',
15    'Mystery': '#DC2026',
16    'Music': '#DC2026',
17    'Fantasy': '#DC2026',
18    'Sci-Fi': '#DC2026',
19    'Animation': '#DC2026',
20    'Sport': '#DC2026',
21    'News': '#DC2026',
22    'Musical': '#DC2026',
23    'War': '#DC2026',
24    'Western': '#DC2026',
25    'Reality-TV': '#DC2026',
26    'Talk-Show': '#DC2026',
27    'Adult': '#DC2026',
28    'Short': '#DC2026',
29    'Game-Show': '#DC2026',
30 }
31
32 # Create a new column 'color' in the dataframe based on genre
33 genre_counts['color'] = genre_counts['genre'].apply(lambda x: color_codes.get(x, '#FCC252'))
34
35 # Define treemap trace
36 treemap_trace = go.Treemap(
37     labels=genre_counts['genre'],
38     parents=[''] * len(genre_counts['genre']),
39     values=genre_counts['count'],
40     marker=dict(
41         colors=genre_counts['color'],
42         line=dict(
43             color=['#000000' if genre in ['Action', 'Adventure', 'Comedy'] else 'rgba(0,0,0,0)' for gen
44             width=2
```

```
45     )
46   )
47 )
48 )
49
50 # Define symbols for specific genres
51 symbols = ['*' if genre in ['Action', 'Adventure', 'Comedy'] else '' for genre in genre_counts['genre']]
52
53 # Add annotations for symbols
54 annotations = [dict(text=symbol, x=genre, y=1, font=dict(size=12)) for genre, symbol in zip(genre_counts['genre'],
55
56 # Create layout
57 layout = go.Layout(
58     title='Tree Map: Movie Production Count by Genre',
59     annotations=annotations,
60     height=800,
61     width=1000,
62     font=dict(
63         family="Times New Roman, serif",
64         size=14,
65         color="black"
66     )
67 )
68
69 # Create figure
70 fig = go.Figure(data=[treemap_trace], layout=layout)
71
72 # Show plot
73 fig.show()
74
```

Tree Map: Movie Production Count by Genre



In a Tree Map the size of the box correlates to the number of movies produced. The top three most profitable genres of Action, Adventure, and Comedy are outlined in black. More commonly made movies are not always correlated to the top grossing genres. This could be due to the overall cost of making Action, Adventure, and Comedy movies which could be determined upon further investigation.

In [70]:

```
1 # Group by 'genre' and sum the 'cpi_avg_domestic_gross' column
2 genre_summary = genre_analysis_explode.groupby('genre')['cpi_avg_domestic_gross'].sum().reset_index()
3
4 # Sort the summary by 'cpi_avg_domestic_gross' in descending order
5 genre_summary = genre_summary.sort_values('cpi_avg_domestic_gross', ascending=False)
6
7 # Display the summary
8 display(genre_summary)
```

| | genre | cpi_avg_domestic_gross |
|----|-------------|------------------------|
| 2 | Adventure | 6.058650e+10 |
| 0 | Action | 5.454206e+10 |
| 5 | Comedy | 4.331613e+10 |
| 8 | Drama | 3.406927e+10 |
| 20 | Sci-Fi | 2.179418e+10 |
| 3 | Animation | 1.899803e+10 |
| 24 | Thriller | 1.671366e+10 |
| 10 | Fantasy | 1.451499e+10 |
| 6 | Crime | 1.235779e+10 |
| 9 | Family | 9.881030e+09 |
| 19 | Romance | 9.533938e+09 |
| 13 | Horror | 7.831989e+09 |
| 4 | Biography | 7.396756e+09 |
| 16 | Mystery | 6.891842e+09 |
| 14 | Music | 2.541608e+09 |
| 12 | History | 2.417894e+09 |
| 22 | Sport | 1.536252e+09 |
| 15 | Musical | 1.111355e+09 |
| 7 | Documentary | 8.277858e+08 |
| 26 | Western | 8.044687e+08 |
| 25 | War | 5.801241e+08 |
| 17 | News | 1.726527e+04 |
| 1 | Adult | 0.000000e+00 |
| 18 | Reality-TV | 0.000000e+00 |
| 11 | Game-Show | 0.000000e+00 |
| 21 | Short | 0.000000e+00 |

| | genre | cpi_avg_domestic_gross |
|----|-----------|------------------------|
| 23 | Talk-Show | 0.000000e+00 |

In [71]:

```
1 # View numeric and word format of the domestic gross values by genre
2
3 # Set the locale for formatting numbers
4 locale.setlocale(locale.LC_ALL, 'en_US.UTF-8')
5
6 # Function to format numbers in both numeric and word format
7 def format_numbers(num):
8     """
9         Function to format numbers in both numeric and word format.
10
11     Input:
12         num : Numeric value to be formatted
13
14     Output:
15         num_formatted : Formatted numeric value with grouping
16         num_words : Numeric value converted to words
17     """
18
19     # Format number in numeric format
20     num_formatted = locale.format_string('%.0f', num, grouping=True)
21
22     # Convert number to words
23     p = inflect.engine()
24     num_words = p.number_to_words(int(num)).replace('-', ' ')
25
26     return num_formatted, num_words
27
28 # Apply the function to the 'cpi_avg_domestic_gross' column
29 genre_summary['numeric_format'], genre_summary['word_format'] = zip(*genre_summary['cpi_avg_domestic_gross'].apply(format_numbers))
30
31 # Display the updated summary without genres that report zero gross
32 display(genre_summary[genre_summary['cpi_avg_domestic_gross'] != 0])
```

| | genre | cpi_avg_domestic_gross | numeric_format | word_format |
|----|-------------|------------------------|----------------|---|
| 2 | Adventure | 6.058650e+10 | 60,586,497,706 | sixty billion, five hundred and eighty six mil... |
| 0 | Action | 5.454206e+10 | 54,542,059,389 | fifty four billion, five hundred and forty two... |
| 5 | Comedy | 4.331613e+10 | 43,316,127,595 | forty three billion, three hundred and sixteen... |
| 8 | Drama | 3.406927e+10 | 34,069,269,087 | thirty four billion, sixty nine million, two h... |
| 20 | Sci-Fi | 2.179418e+10 | 21,794,184,801 | twenty one billion, seven hundred and ninety f... |
| 3 | Animation | 1.899803e+10 | 18,998,033,334 | eighteen billion, nine hundred and ninety eigh... |
| 24 | Thriller | 1.671366e+10 | 16,713,659,373 | sixteen billion, seven hundred and thirteen mi... |
| 10 | Fantasy | 1.451499e+10 | 14,514,988,614 | fourteen billion, five hundred and fourteen mi... |
| 6 | Crime | 1.235779e+10 | 12,357,785,265 | twelve billion, three hundred and fifty seven ... |
| 9 | Family | 9.881030e+09 | 9,881,029,880 | nine billion, eight hundred and eighty one mil... |
| 19 | Romance | 9.533938e+09 | 9,533,937,567 | nine billion, five hundred and thirty three mi... |
| 13 | Horror | 7.831989e+09 | 7,831,988,544 | seven billion, eight hundred and thirty one mi... |
| 4 | Biography | 7.396756e+09 | 7,396,755,647 | seven billion, three hundred and ninety six mi... |
| 16 | Mystery | 6.891842e+09 | 6,891,842,100 | six billion, eight hundred and ninety one mill... |
| 14 | Music | 2.541608e+09 | 2,541,608,385 | two billion, five hundred and forty one millio... |
| 12 | History | 2.417894e+09 | 2,417,893,750 | two billion, four hundred and seventeen millio... |
| 22 | Sport | 1.536252e+09 | 1,536,252,159 | one billion, five hundred and thirty six milli... |
| 15 | Musical | 1.111355e+09 | 1,111,355,064 | one billion, one hundred and eleven million, t... |
| 7 | Documentary | 8.277858e+08 | 827,785,784 | eight hundred and twenty seven million, seven ... |
| 26 | Western | 8.044687e+08 | 804,468,702 | eight hundred and four million, four hundred a... |
| 25 | War | 5.801241e+08 | 580,124,132 | five hundred and eighty million, one hundred a... |
| 17 | News | 1.726527e+04 | 17,265 | seventeen thousand, two hundred and sixty five |

In [72]:

```
1 # Calculate the total cpi_avg_domestic_gross
2 total_cpi_avg_domestic_gross = genre_summary['cpi_avg_domestic_gross'].sum()
3
4 # Calculate the percentage contribution for each genre
5 genre_summary['percentage_contribution'] = (genre_summary['cpi_avg_domestic_gross']) / total_cpi_avg_dor
6
7 # Display the table
8 percentage_table = genre_summary[['genre', 'percentage_contribution']].sort_values(by='percentage_conti
9
10 # Display the updated DataFrame
11 print('Total number of results:', len(genre_analysis_explode))
12 display(percentage_table)
```

Total number of results: 249658

| | genre | percentage_contribution |
|----|-------------|-------------------------|
| 2 | Adventure | 18.457556 |
| 0 | Action | 16.616130 |
| 5 | Comedy | 13.196172 |
| 8 | Drama | 10.379135 |
| 20 | Sci-Fi | 6.639555 |
| 3 | Animation | 5.787713 |
| 24 | Thriller | 5.091783 |
| 10 | Fantasy | 4.421962 |
| 6 | Crime | 3.764775 |
| 9 | Family | 3.010236 |
| 19 | Romance | 2.904495 |
| 13 | Horror | 2.386000 |
| 4 | Biography | 2.253407 |
| 16 | Mystery | 2.099586 |
| 14 | Music | 0.774296 |
| 12 | History | 0.736607 |
| 22 | Sport | 0.468016 |
| 15 | Musical | 0.338572 |
| 7 | Documentary | 0.252183 |
| 26 | Western | 0.245080 |
| 25 | War | 0.176734 |
| 17 | News | 0.000005 |
| 1 | Adult | 0.000000 |
| 18 | Reality-TV | 0.000000 |
| 11 | Game-Show | 0.000000 |
| 21 | Short | 0.000000 |

| | genre | percentage_contribution |
|----|-----------|-------------------------|
| 23 | Talk-Show | 0.000000 |

It is useful to view the genre by percentages that they contribute to that total to get a sense for how much each genre contributes to the overall domestic gross revenue. For example, Adventure movies account for almost 20% of all reported domestic revenue meaning that they contribute substantially to the overall total domestic gross revenue.

→ In this next section, let's explore how much the top 100 most profitable movies of all time contribute to the total domestic gross revenue by genre.

```
In [73]: 1 print('Total number of results:', len(top_profit))
2 display(top_profit.head())
```

Total number of results: 100

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_dome |
|------|--|-------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|--------------|
| 6814 | gone with the wind | | NaN | 1939 | 1939-12-15 | 3900000.0 | 390525192.0 | 198680470.0 | 8.560706e+09 | 4.3% |
| 7293 | snow white and the seven dwarfs | | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.9% |
| 5928 | star wars ep. iv: a new hope | | NaN | 1977 | 1977-05-25 | 11000000.0 | 786598007.0 | 460998007.0 | 3.955082e+09 | 2.3% |
| 7497 | bambi | | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.9% |
| 7073 | pinocchio | | NaN | 1940 | 1940-02-09 | 2289247.0 | 84300000.0 | 84300000.0 | 1.834741e+09 | 1.8% |

```
In [74]: 1 # Inner merge 'top_profit' with 'genre_analysis_explode' on 'title', 'year', and 'genre'
2 top_profit_genre = pd.merge(top_profit, genre_analysis_explode[['title', 'release_year', 'genre']], le
3
4 # Display the resulting DataFrame
5 print('Total number of results:', len(top_profit_genre))
6 display(top_profit_genre.head())
7
```

Total number of results: 128

| | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domestic |
|---|--|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|------------------|
| 0 | gone with the wind | NaN | 1939 | 1939-12-15 | 3900000.0 | 390525192.0 | 198680470.0 | 8.560706e+09 | 4.35521 |
| 1 | snow white and the seven dwarfs | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.91299 |
| 2 | star wars ep. iv: a new hope | NaN | 1977 | 1977-05-25 | 11000000.0 | 786598007.0 | 460998007.0 | 3.955082e+09 | 2.31791 |
| 3 | bambi | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.92162 |
| 4 | pinocchio | NaN | 1940 | 1940-02-09 | 2289247.0 | 84300000.0 | 84300000.0 | 1.834741e+09 | 1.83474 |

```
In [75]: 1 # Why are there more than 100 entries in the top_profit_genre df?
2 # Consider unique titles
3 unique_titles_count = top_profit_genre['title'].nunique()
4 unique_titles_count
```

Out[75]: 99

In [76]:

```
1 # Consider duplicate titles
2 duplicate_titles = top_profit[top_profit['title'].duplicated(keep=False)]
3 display(duplicate_titles)
```

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_domestic |
|------|--|-------------------------------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|------------------|
| 5165 | | beauty and the beast | NaN | 1991 | 1991-11-13 | 20000000.0 | 6.084311e+08 | 376057266.0 | 1.361161e+09 | 8.4130 |
| 3439 | | beauty and the beast | NaN | 2017 | 2017-03-17 | 160000000.0 | 1.259200e+09 | 504014165.0 | 1.565277e+09 | 6.2652 |

This shows 99 unique titles in the top_profit_genre dataframe because some titles such as Beauty and the Beast (release years: 1991 and 2017) were top profit movies in multiple release years. This and the exploded genre count account for the 128 values shown in the top_profit_genre dataframe.

In [77]:

```
1 # Calculate the total number of movies
2 total_movies = top_profit_genre.shape[0]
3
4 # Calculate the percentage of movies for each genre
5 genre_percentage = top_profit_genre['genre'].value_counts(normalize=True) * 100
6
7 # Display the resulting table
8 print('Total number of results:', len(genre_percentage))
9 display(genre_percentage)
```

Total number of results: 10

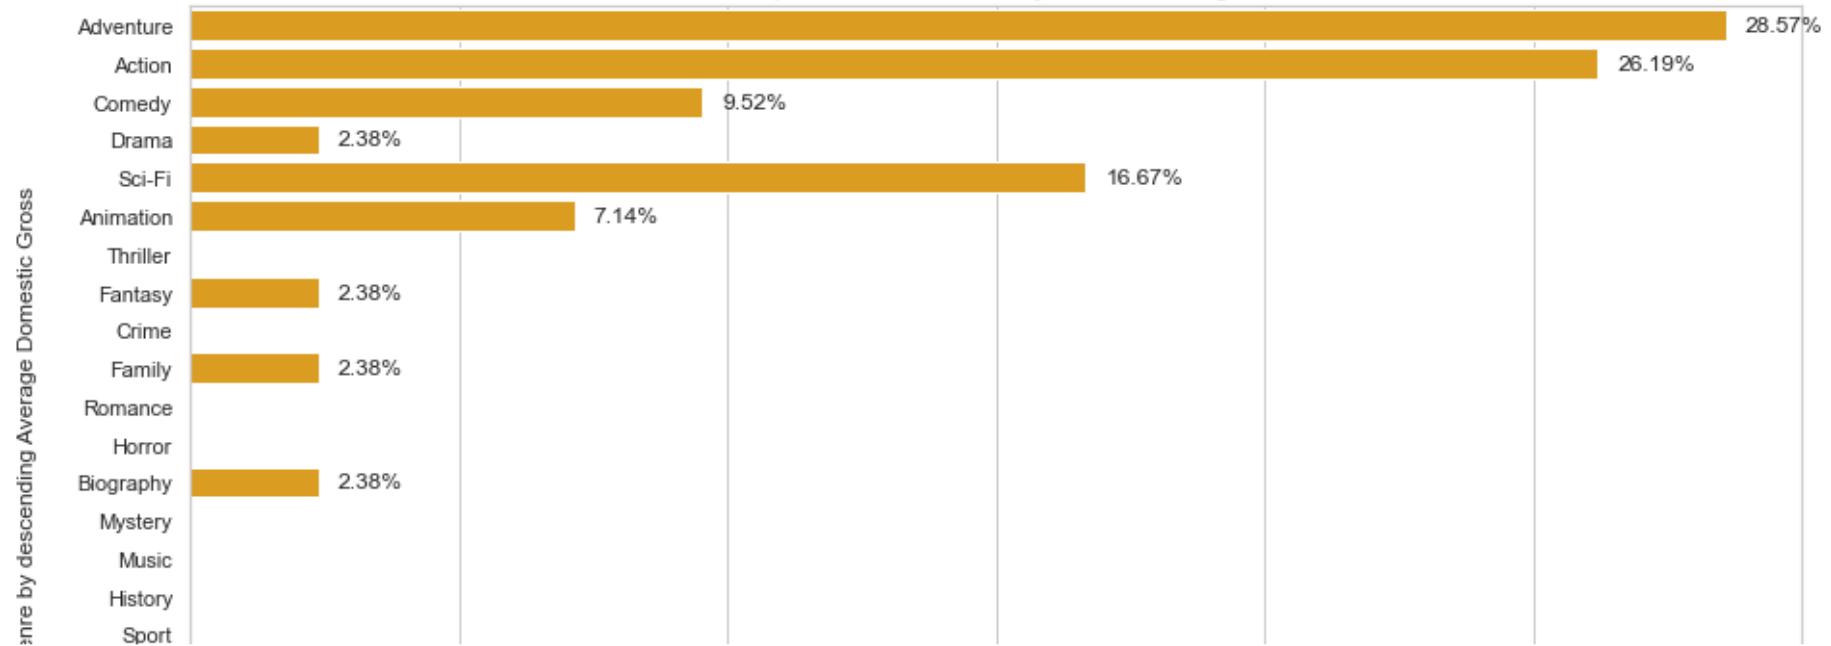
| | |
|-----------|-----------|
| Adventure | 28.571429 |
| Action | 26.190476 |
| Sci-Fi | 16.666667 |
| Comedy | 9.523810 |
| Animation | 7.142857 |
| Drama | 2.380952 |
| Family | 2.380952 |
| Biography | 2.380952 |
| Fantasy | 2.380952 |
| Musical | 2.380952 |

Name: genre, dtype: float64

In [78]:

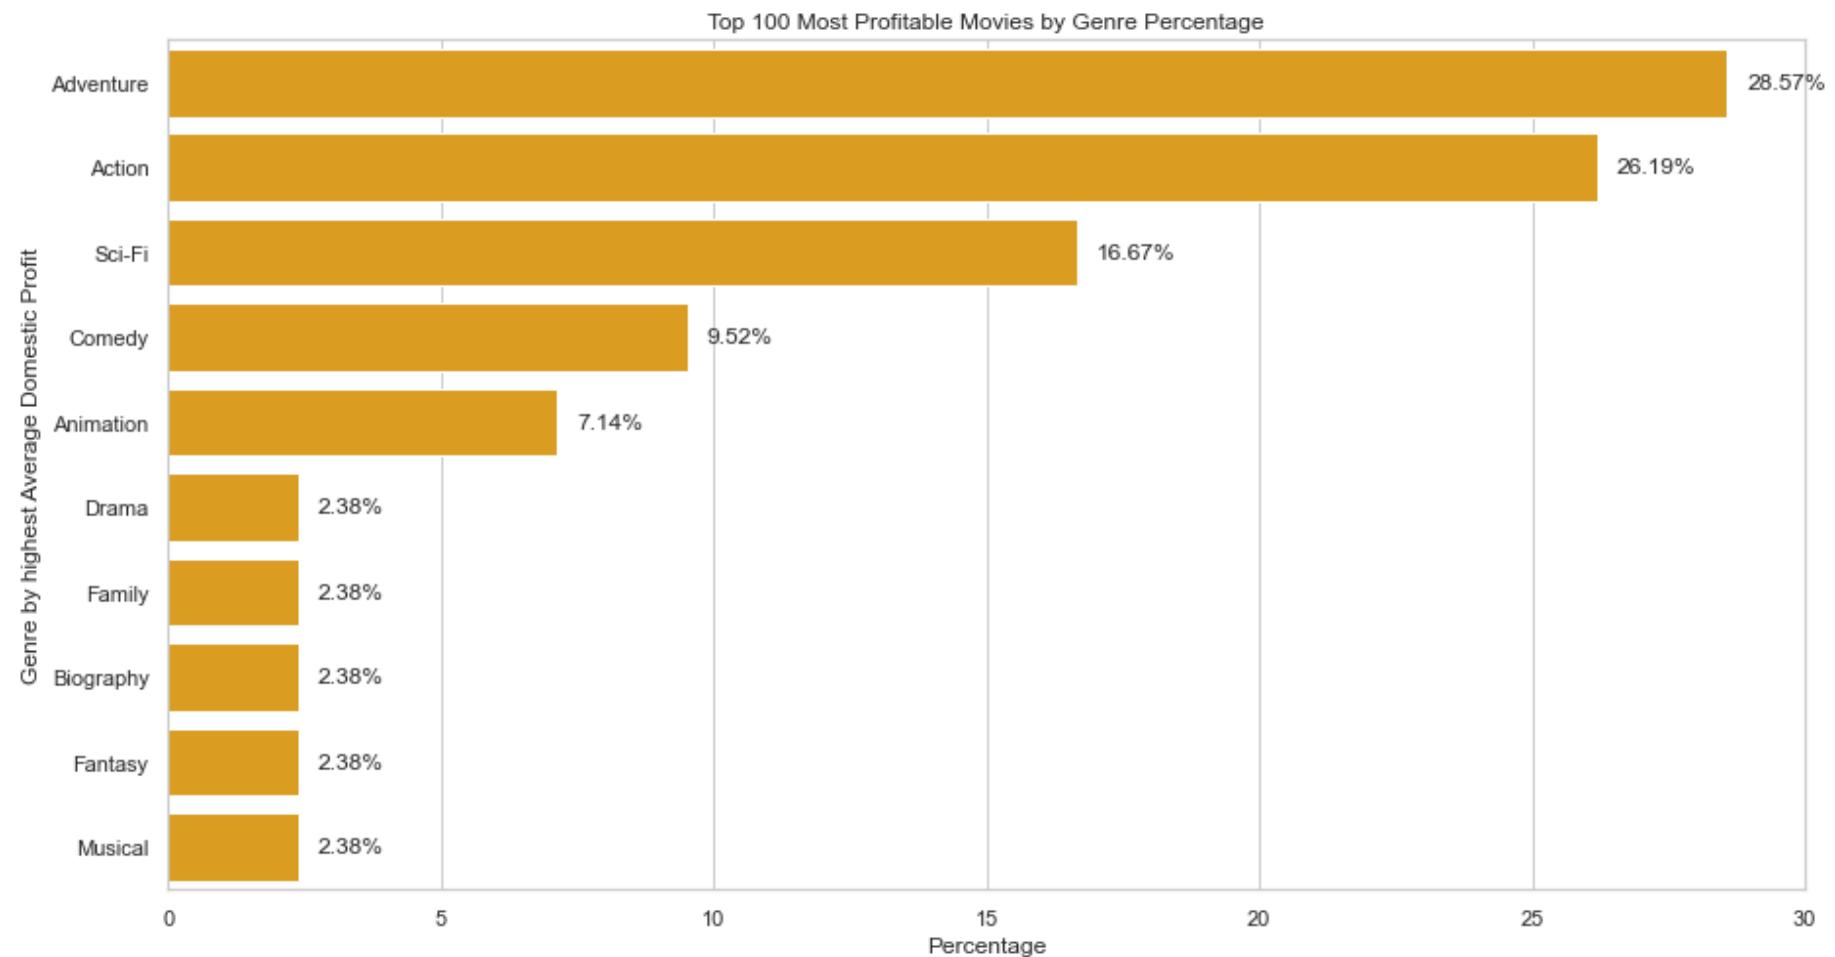
```
1 # Convert the genre_percentage Series to a DataFrame
2 genre_percentage_df = genre_percentage.reset_index()
3 genre_percentage_df.columns = ['genre', 'percentage']
4
5
6 # Define the order of genres
7 genre_order = ['Adventure', 'Action', 'Comedy', 'Drama', 'Sci-Fi', 'Animation', 'Thriller', 'Fantasy',
8                 'Family', 'Romance', 'Horror', 'Biography', 'Mystery', 'Music', 'History', 'Sport', 'Mus
9                 'Documentary', 'Western', 'War', 'News']
10
11 # Set the style
12 sns.set(style="whitegrid")
13
14 # Create a horizontal bar graph with specified order
15 plt.figure(figsize=(15, 8))
16 bar_plot = sns.barplot(x='percentage', y='genre', data=genre_percentage_df, order=genre_order, color='r
17
18 # Add percentages on the right of each bar
19 for p in bar_plot.patches:
20     bar_plot.annotate(f'{p.get_width():.2f}%', (p.get_width(), p.get_y() + p.get_height() / 2.),
21                       ha='left', va='center', xytext=(10, 0), textcoords='offset points')
22
23 # Set plot title and labels
24 plt.title('Top 100 Profitable Movies by Genre Percentage')
25 plt.xlabel('Percentage')
26 plt.ylabel('Genre by descending Average Domestic Gross')
27
28 # Show the plot
29 plt.show()
```

Top 100 Profitable Movies by Genre Percentage



In [79]:

```
1 # Set the style
2 sns.set(style="whitegrid")
3
4 # Create a horizontal bar graph
5 plt.figure(figsize=(15, 8))
6 bar_plot = sns.barplot(x='percentage', y='genre', data=genre_percentage_df, color="#F9A602")
7
8 # Add percentages on the right of each bar
9 for p in bar_plot.patches:
10     bar_plot.annotate(f'{p.get_width():.2f}%', (p.get_width(), p.get_y() + p.get_height() / 2.),
11                        ha='left', va='center', xytext=(10, 0), textcoords='offset points')
12
13 # Set plot title and labels
14 plt.title('Top 100 Most Profitable Movies by Genre Percentage')
15 plt.xlabel('Percentage')
16 plt.ylabel('Genre by highest Average Domestic Profit')
17
18 # Show the plot
19 plt.show()
```

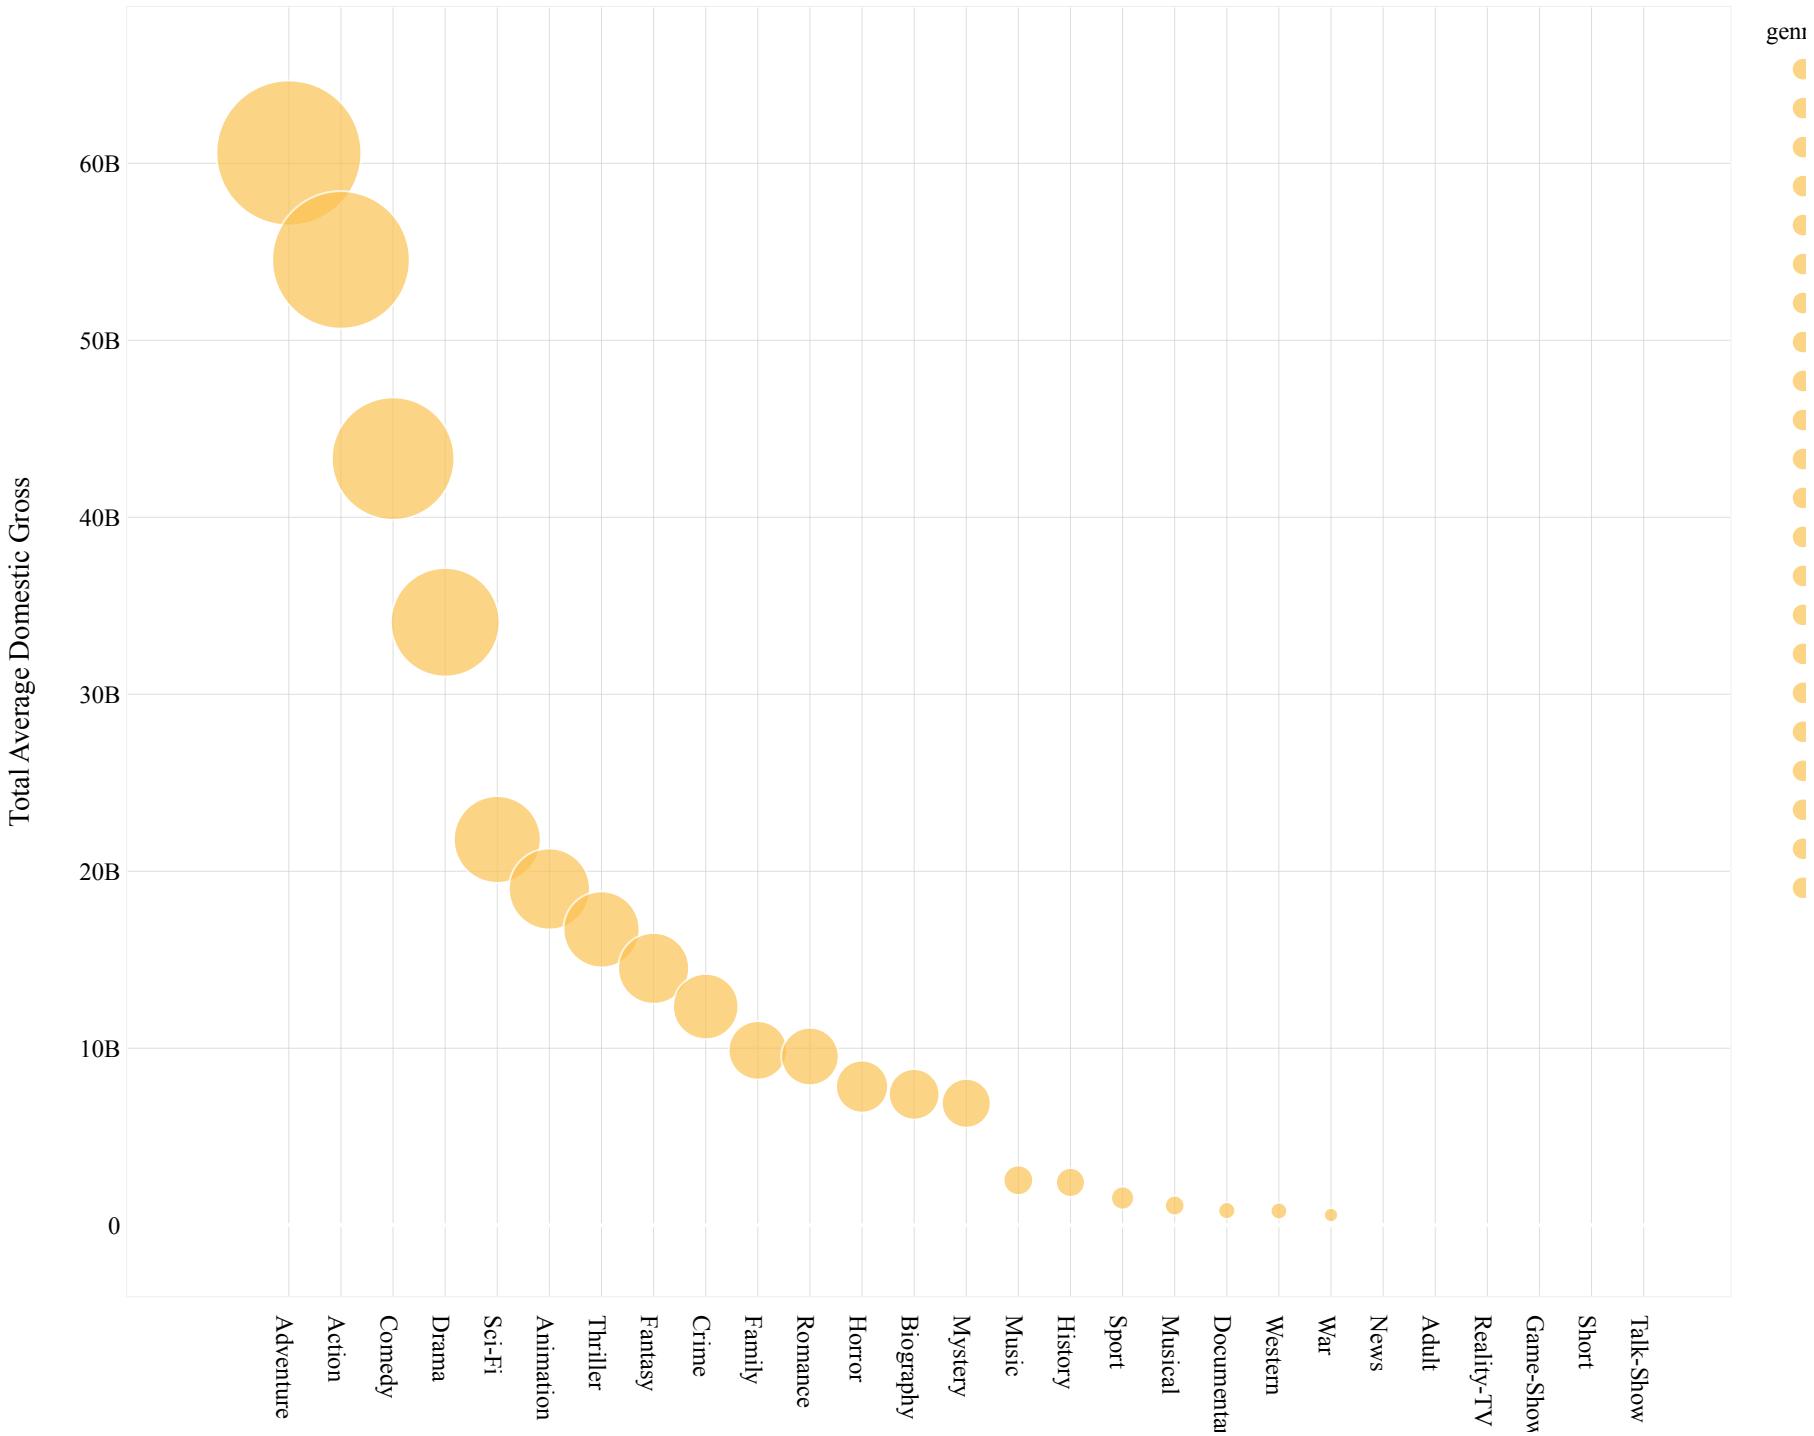


About 64% of top 100 most profitable movies of all time fall into the top three, all-time grossing genres of Action, Adventure, and Comedy.

In [80]:

```
1 # Define the custom colors
2 background_color = '#FFFFFF' # White background
3 grid_line_color = '#d3d3d3'
4 bubble_color = '#FCC252'
5 text_color = '#000000'
6 outline_color = '#eaeaea'
7
8 # Creating the bubble plot using Plotly with custom colors
9 fig = px.scatter(genre_summary, x='genre', y='cpi_avg_domestic_gross', size='cpi_avg_domestic_gross',
10                  hover_name='genre', size_max=50, color='genre', color_discrete_sequence=[bubble_color])
11
12 # Set plot characteristics
13 fig.update_layout(title='Bubble Plot: Total Average Domestic Gross by Genre',
14                     xaxis_title='Genre',
15                     yaxis_title='Total Average Domestic Gross',
16                     plot_bgcolor=background_color,
17                     paper_bgcolor=background_color,
18                     xaxis=dict(showgrid=True, gridcolor=grid_line_color, linecolor=outline_color),
19                     yaxis=dict(showgrid=True, gridcolor=grid_line_color, linecolor=outline_color),
20                     font=dict(family="Times New Roman", color=text_color),
21                     height=800,
22                     width=1000)
23
24 # Add borders around the plot
25 fig.update_xaxes(showline=True, linewidth=1, linecolor=outline_color, mirror=True)
26 fig.update_yaxes(showline=True, linewidth=1, linecolor=outline_color, mirror=True)
27
28 # Show plot
29 fig.show()
30
```

Bubble Plot: Total Average Domestic Gross by Genre



Genre Analysis Summary

The analysis of the most profitable movie genres reveals trends into the domestic gross earnings. Adventure emerges as the most financially successful genre, followed by Action and Comedy. Some genres, such as Talk-Show, Short, Reality-TV, Adult, and Game-Show, show zero values in the adjusted average domestic gross. This can be attributed to the integration of movie genre data from one dataset and financial data from another, leading to instances of missing financial information for specific genres.

Considering the top 100 most profitable movies, 64% of titles fall into the overall top grossing genres of Action, Adventure, and Comedy.

Additionally, it's essential to recognize the potential impact of mixed-genre films on the dataset. Some movies may belong to more than one genre, introducing complexity and potential ambiguity in attributing financial success solely to a specific genre.

Considering these limitations, future analyses could explore the correlation between mixed-genre films and financial performance.

Release Date Analysis

The month that a movie is released may greatly influence profitability so when is the most profitable time to release a movie? By investigating the release month of a movie and referencing its adjusted average domestic gross, it can be determined if there is a more profitable time to release movies.

In [81]:

```
1 # Create a copy of the DataFrame to avoid altering the original
2 release_date_analysis = main_df_clean.copy()
3
4 # Select specific columns to create a new DataFrame
5 selected_columns = ['movie_id', 'release_year', 'title', 'release_date', 'cpi_avg_domestic_gross']
6
7 # Create a new DataFrame with only the selected columns
8 release_date_summary = release_date_analysis[selected_columns].copy()
9
10 # Create a release_month column
11 release_date_summary['release_month']=pd.DatetimeIndex(release_date_summary['release_date']).month
12
13 # Display the new DataFrame
14 print('Total number of results:', len(release_date_summary))
15 display(release_date_summary.head())
```

Total number of results: 161521

| | movie_id | release_year | title | release_date | cpi_avg_domestic_gross | release_month |
|---|-----------|--------------|---------------------------------|--------------|------------------------|---------------|
| 0 | tt0063540 | 2013.0 | sunghursh | NaN | NaN | NaN |
| 1 | tt0066787 | 2019.0 | one day before the rainy season | NaN | NaN | NaN |
| 2 | tt0069204 | 2018.0 | sabse bada sukh | NaN | NaN | NaN |
| 3 | tt0100275 | 2017.0 | the wandering soap opera | NaN | NaN | NaN |
| 4 | tt0111414 | 2018.0 | a thin life | NaN | NaN | NaN |

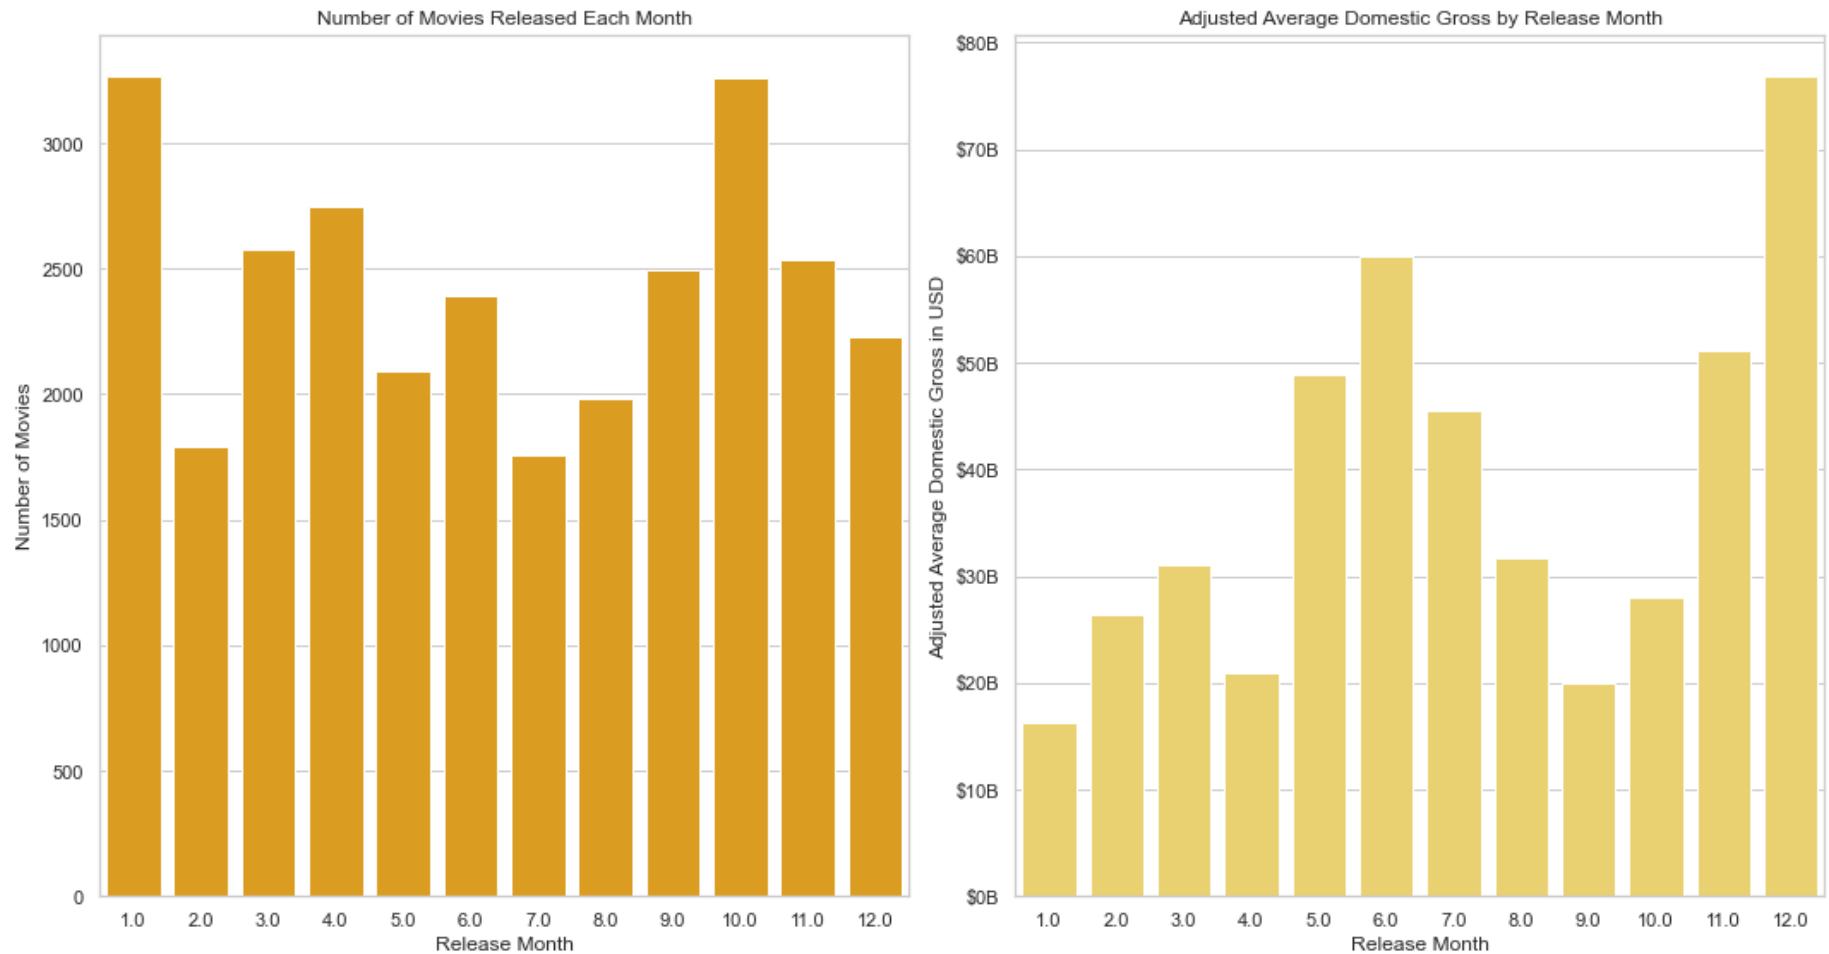
This function will be useful to create subsequent multiple visualizations on this data. Visualizing the number of movies released each month and the adjusted average domestic gross for each month will provide insight on the best range of dates to release a movie.

In [82]:

```
1 def billions_formatter(x, pos):
2     """
3         Function to format y-axis labels to display in billions.
4
5         Input:
6         x : Value to be formatted
7         pos : Position of the tick on the axis
8
9         Output:
10        formatted_string : Formatted string representing the value in billions
11        """
12        return f'{x/1e9:.0f}B'
13
14 def release_month_visual(df):
15     """
16         Visualizes movie release patterns based on months.
17
18         Input:
19         df : Pandas dataframe containing information about movies
20
21         Output:
22         None (Displays a matplotlib figure)
23         """
24 # Create a figure with 1 row and 2 columns
25 fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(15, 8))
26 # Create a count plot for the number of movies released each month
27 sns.countplot(x='release_month', color='#F9A602', data=df, ax=ax1)
28 ax1.set_xlabel('Release Month')
29 ax1.set_ylabel('Number of Movies')
30 ax1.set_title('Number of Movies Released Each Month')
31 # Create a bar plot for the avg domestic gross by release month
32 sns.barplot(x='release_month', y='cpi_avg_domestic_gross', estimator=sum, ci=None, color="#FDCC5C")
33 ax2.set_xlabel('Release Month')
34 ax2.set_ylabel('Adjusted Average Domestic Gross in USD')
35 ax2.set_title('Adjusted Average Domestic Gross by Release Month')
36
37 # Adjusting the formatting of the y-axis for the second plot
38 formatter = FuncFormatter(billions_formatter)
39 ax2.yaxis.set_major_formatter(formatter)
40
41 plt.tight_layout()
42 # Show the plot
```

```
43 plt.show()
```

```
In [83]: 1 # Run the function on release_date_summary  
2 release_month_visual(release_date_summary)
```



It is shown here that the most popular months for releasing movies are January, March, April, October, and November. The months where the average domestic gross is highest are May, June, July, November, and December.

```
In [84]: 1 # View top_profit dataframe
2 print('Total number of results:', len(top_profit))
3 display(top_profit.head())
```

Total number of results: 100

| | | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg_domestic_gross | cpi_worldwide_gross | cpi_avg_dome |
|------|--|-------|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|--------------|
| 6814 | gone with the wind | | NaN | 1939 | 1939-12-15 | 3900000.0 | 390525192.0 | 198680470.0 | 8.560706e+09 | 4.3% |
| 7293 | snow white and the seven dwarfs | | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.9% |
| 5928 | star wars ep. iv: a new hope | | NaN | 1977 | 1977-05-25 | 11000000.0 | 786598007.0 | 460998007.0 | 3.955082e+09 | 2.3% |
| 7497 | bambi | | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.9% |
| 7073 | pinocchio | | NaN | 1940 | 1940-02-09 | 2289247.0 | 84300000.0 | 84300000.0 | 1.834741e+09 | 1.8% |

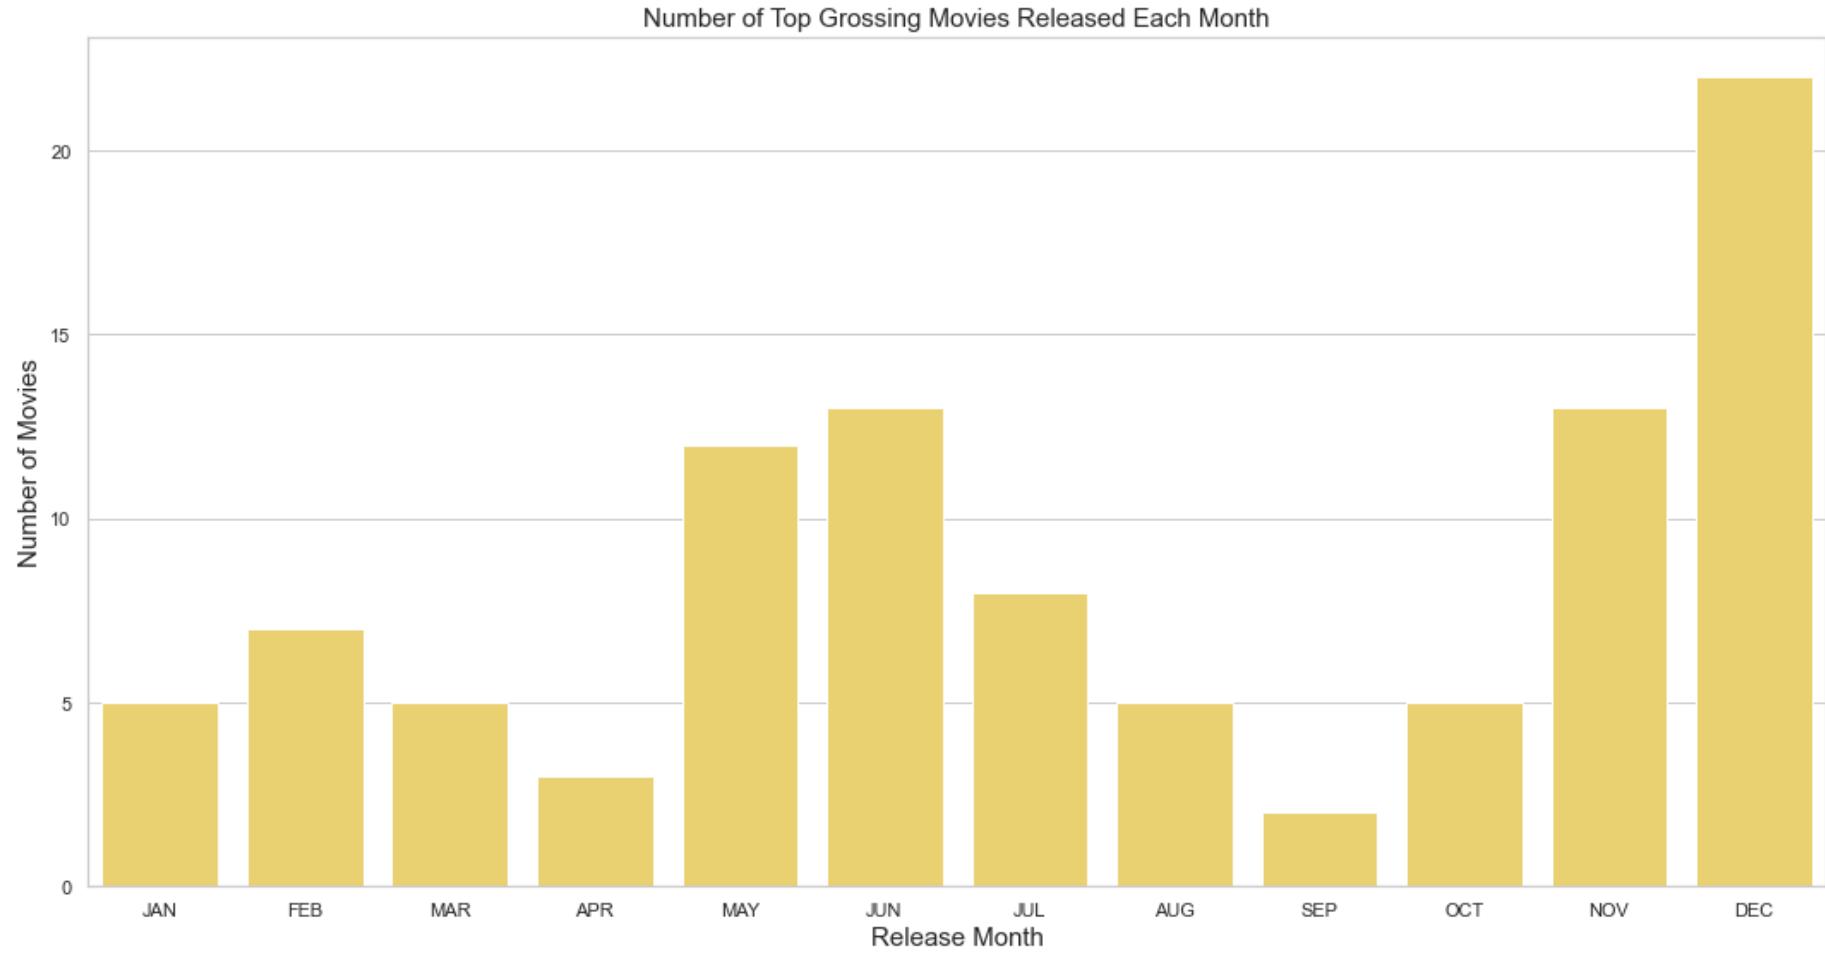
```
In [85]: 1 # Create a release_month column
2 top_profit['release_month']=pd.DatetimeIndex(top_profit['release_date']).month
```

In [86]:

```
1 def release_month_visual_tg(df):
2     """
3         Visualizes the number of top-grossing movies released each month.
4
5     Input:
6         df : Pandas dataframe containing information about movies
7
8     Output:
9         None (Displays a matplotlib figure)
10    """
11    # Create a Counter of months from the available data
12    months_counter = Counter(df['release_month'])
13    # Define all months you want to display
14    all_months = ['JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC']
15
16    # Create counts for all months, filling missing months with zero counts
17    counts_per_month = {month: months_counter.get(index + 1, 0) for index, month in enumerate(all_months)}
18    plt.figure(figsize=(15, 8))
19    ax = sns.barplot(x=list(counts_per_month.keys()), y=list(counts_per_month.values()),
20                      color="#FDCC5C", label='Number of Movies', ci=None)
21
22    # Set positions and labels for x-axis ticks
23    ax.set_xticks(range(len(all_months)))
24    ax.set_xticklabels(all_months)
25
26    ax.set_xlabel('Release Month', fontfamily='Arial', fontsize=15)
27    ax.set_ylabel('Number of Movies', fontfamily='Arial', fontsize=15)
28    ax.set_title('Number of Top Grossing Movies Released Each Month', fontfamily='Arial', fontsize=15)
29
30
31    plt.tight_layout()
32
33    # Show the plot
34    plt.show()
```

In [87]:

```
1 # Run the function on top_profit  
2 release_month_visual_tg(top_profit)
```



This visual somewhat agrees with what was shown previously in that **May, June, July, and November/December are the best months to release movies.**

```
In [88]: 1 # Format 'release_date' in the original DataFrame
2 release_date_summary['release_date'] = pd.to_datetime(release_date_summary['release_date'], errors='coerce')
3
4 # Create a new column 'day_of_week' in the original DataFrame
5 release_date_summary['day_of_week'] = release_date_summary['release_date'].dt.day_name()
```

```
In [89]: 1 # Count occurrences of each day of the week
2 day_counts = release_date_summary['day_of_week'].value_counts()
3
4 # Find the most popular day
5 most_popular_day = day_counts.idxmax()
6
7 # Display the results
8 print("Counts of each day of the week:")
9 print(day_counts)
10
11 print("\nThe most popular day to release a movie is:", most_popular_day)
12
```

Counts of each day of the week:

| | |
|-----------|-------|
| Friday | 11585 |
| Tuesday | 4366 |
| Saturday | 3483 |
| Thursday | 2664 |
| Wednesday | 2589 |
| Sunday | 2495 |
| Monday | 1949 |

Name: day_of_week, dtype: int64

The most popular day to release a movie is: Friday

This list shows that the most popular day of the week to release a movie is Friday.

Release Date Summary

This analysis highlights a relationship between the release date and the success of top-grossing movies. Caution must be used here because additional factors may influence a movie's overall profit such as the economy, viewing platform trends, distribution strategy (wide vs limited release), advertising budget, any award recognition, and the movie's enduring appeal over time.

A comprehensive analysis that incorporates specific factors expressedly communicated by the film company as priorities based on the company staff strengths is essential for a nuanced understanding of the dynamics influencing the financial success of its movies is needed to

Runtime Analysis

Runtime can impact the type of audience that views the movie. Kid and family films may tend to be shorter while action movies may be longer. Depending on the intended audience, runtime may impact the financial success of the film.

In [90]:

```
1 # Create a copy of the DataFrame to avoid altering the original
2 runtime_analysis = main_df_clean.copy()
3
4 # Select specific columns to create a new DataFrame
5 selected_columns = ['title', 'runtime_minutes', 'cpi_avg_domestic_gross', 'genres', 'release_year']
6
7 # Create a new DataFrame with only the selected columns
8 runtime_summary = runtime_analysis[selected_columns].copy()
9
10 # Display the new DataFrame
11 runtime_summary.head()
```

Out[90]:

| | title | runtime_minutes | cpi_avg_domestic_gross | genres | release_year |
|---|---------------------------------|-----------------|------------------------|----------------------|--------------|
| 0 | sunghursh | 175.0 | NaN | Action,Crime,Drama | 2013.0 |
| 1 | one day before the rainy season | 114.0 | NaN | Biography,Drama | 2019.0 |
| 2 | sabse bada sukh | NaN | NaN | Comedy,Drama | 2018.0 |
| 3 | the wandering soap opera | 80.0 | NaN | Comedy,Drama,Fantasy | 2017.0 |
| 4 | a thin life | 75.0 | NaN | Comedy | 2018.0 |

```
In [91]: 1 # Get statistics for movie runtime  
2 runtime_summary[['title', 'runtime_minutes']].describe()
```

Out[91]:

| | runtime_minutes |
|-------|-----------------|
| count | 113046.000000 |
| mean | 86.221193 |
| std | 167.322817 |
| min | 1.000000 |
| 25% | 70.000000 |
| 50% | 87.000000 |
| 75% | 99.000000 |
| max | 51420.000000 |

→ There appears to be an outlier because the standard deviation is huge. There potentially aren't too many outliers because the mean and median are close together, so let's take a closer look.

```
In [92]: 1 # Order the runtime minutes by longest to shortest  
2 print('Total number of results:', len(runtime_summary.sort_values(by='runtime_minutes', ascending=False)))  
3 display((runtime_summary.sort_values(by='runtime_minutes', ascending=False)).head())
```

Total number of results: 161521

| | title | runtime_minutes | cpi_avg_domestic_gross | genres | release_year |
|--------|----------------------|-----------------|------------------------|----------------------|--------------|
| 116541 | logistics | 51420.0 | NaN | Documentary | 2012.0 |
| 37951 | modern times forever | 14400.0 | NaN | Documentary | 2011.0 |
| 108172 | nari | 6017.0 | NaN | Documentary | 2017.0 |
| 75288 | hunger! | 6000.0 | NaN | Documentary,Drama | 2015.0 |
| 76594 | london ec1 | 5460.0 | NaN | Comedy,Drama,Mystery | 2015.0 |

In [93]:

```

1 # Make a variable to take a closer look at the outliers
2 sorted_runtime = runtime_summary.sort_values(by='runtime_minutes', ascending=False)
3
4 # Run a conditional statement to view some of the outliers
5 print(len(sorted_runtime[sorted_runtime['runtime_minutes']>180]))
6 display((sorted_runtime[sorted_runtime['runtime_minutes']>180]).head())

```

532

| | | title | runtime_minutes | cpi_avg_domestic_gross | genres | release_year |
|--------|----------------------|------------|-----------------|------------------------|----------------------|--------------|
| 116541 | | logistics | 51420.0 | NaN | Documentary | 2012.0 |
| 37951 | modern times forever | | 14400.0 | NaN | Documentary | 2011.0 |
| 108172 | | nari | 6017.0 | NaN | Documentary | 2017.0 |
| 75288 | | hunger! | 6000.0 | NaN | Documentary,Drama | 2015.0 |
| 76594 | | london ec1 | 5460.0 | NaN | Comedy,Drama,Mystery | 2015.0 |

→ After some internet fact checking, some of the runtime_minutes reported are false. A movie of 300 minutes would mean five hours which is a tremendous undertaking for production. So considering that the mean statistics are skewed by these outliers, using the 95% quantile seems like a reasonable cut off for reliably reported film runtimes.

Additionally, according to the Academy of Motion Picture Arts and Sciences, the American Film Institute and the British Film Institute, a feature film runs for more than 40 minutes.

In [94]:

```

1 # Calculate the 5th and 95th percentiles
2 percentile_5 = runtime_summary['runtime_minutes'].quantile(0.05)
3 percentile_95 = runtime_summary['runtime_minutes'].quantile(0.95)
4
5 # Filter the DataFrame based on the percentiles
6 runtime_summary_inner_quantile = runtime_summary[((runtime_summary['runtime_minutes'] >= percentile_5) &
7

```

```
In [95]: 1 runtime_summary_inner_quantile.describe()
```

Out[95]:

| | runtime_minutes | cpi_avg_domestic_gross | release_year |
|--------------|-----------------|------------------------|---------------|
| count | 101902.000000 | 7.014000e+03 | 150377.000000 |
| mean | 84.811927 | 6.322840e+07 | 2014.144763 |
| std | 19.605914 | 1.419309e+08 | 4.222397 |
| min | 44.000000 | 1.307975e+02 | 1915.000000 |
| 25% | 72.000000 | 6.205125e+05 | 2012.000000 |
| 50% | 87.000000 | 1.768768e+07 | 2015.000000 |
| 75% | 97.000000 | 7.225980e+07 | 2017.000000 |
| max | 130.000000 | 4.355276e+09 | 2115.000000 |

→ Looking at the inner quantile, these values make much more sense.

Comparing the two measurements of spread, the standard deviation and the inter quartile range, before we removed the outliers the inter quartile range was 26 to a standard deviation of 169. After removing the outliers the inter quartile range is 26 to a standard deviation of 20.

This is excellent because for a perfect normal distribution the inter quartile and the standard deviation have a ratio of 1.35.

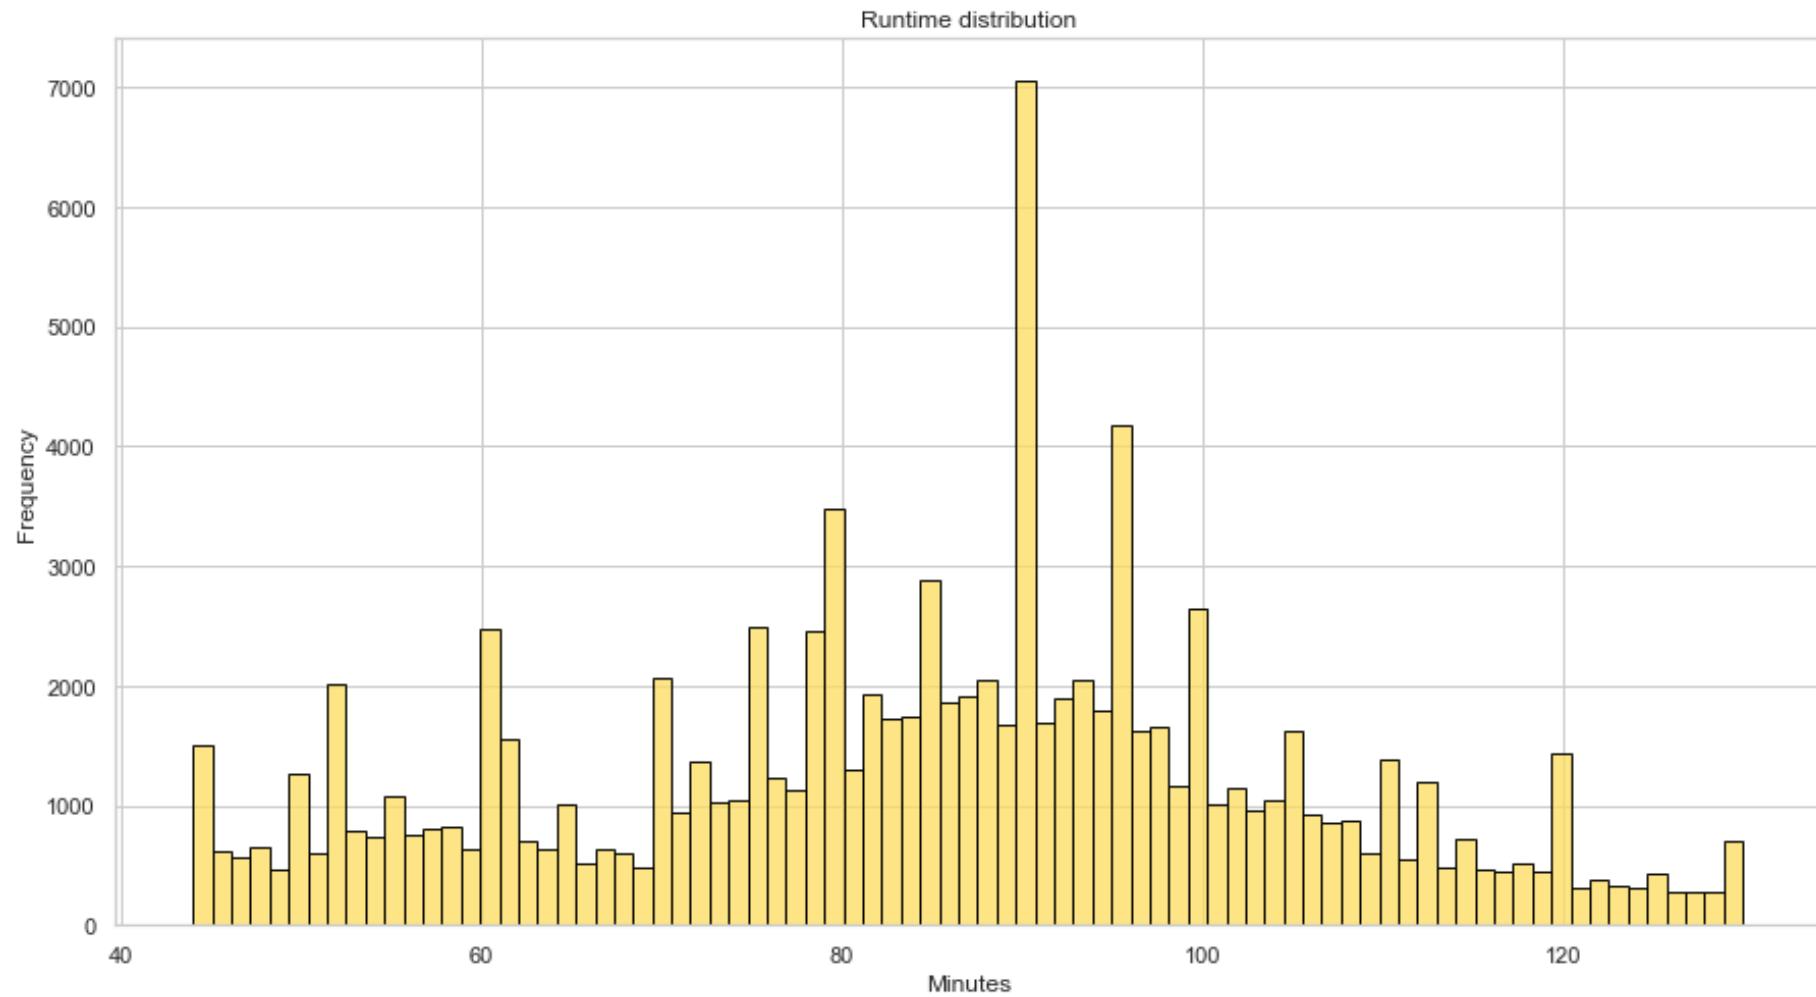
The inter quartile range is 26 to a standard deviation of 20 and the ratio here is 1.3. So after the removal of the outliers this data can confidently be considered normally distributed.

In [96]:

```
1 # Define function to plot
2
3 def runtime_visual(df, title, plot_range=None, bar_color="#FDDC5C", edge_color='black'):
4     """
5         Creates a Seaborn histplot of the 'runtime_minutes' column within a specified range.
6
7         Input:
8             df : Pandas dataframe containing information about movies
9             title: A string representing the title of the plot
10            plot_range: Optional tuple specifying the range of runtimes to plot (e.g., (min_runtime, max_runtime))
11
12        Output:
13            None (Displays a matplotlib figure)
14        """
15    plt.figure(figsize=(15, 8))
16
17    if plot_range:
18        # Filter data within the specified range
19        df_filtered = df[(df['runtime_minutes'] >= plot_range[0]) & (df['runtime_minutes'] <= plot_range[1])]
20        sns.histplot(df_filtered['runtime_minutes'], kde=False, color=bar_color, edgecolor=edge_color)
21        plt.title(f'{title} (Within {plot_range[0]} to {plot_range[1]} Minutes)")
22    else:
23        sns.histplot(df['runtime_minutes'], kde=False, color=bar_color, edgecolor=edge_color)
24        plt.title(title)
25
26    plt.xlabel('Minutes')
27    plt.ylabel('Frequency')
28    return plt.show()
29
```

In [97]:

```
1 # Call runtime_visual function on movies DataFrame  
2 runtime_visual(runtime_summary_inner_quantile, 'Runtime distribution')  
3
```



The median is approximately 87 minutes.

```
In [98]: 1 # Get number of missing values per column  
2 runtime_summary_inner_quantile.isna().sum()
```

```
Out[98]: title          0  
runtime_minutes      48475  
cpi_avg_domestic_gross 143363  
genres            22519  
release_year        0  
dtype: int64
```

```
In [99]: 1 # Percent of missing values for runtime  
2 (runtime_summary_inner_quantile['runtime_minutes'].isna().sum()/len(runtime_summary_inner_quantile)*100)
```

```
Out[99]: 32.24
```

→ This is significant because runtime is essential in this part of the analysis. The missing values will have to be filled.

```
In [100]: 1 # Consider mean  
2 runtime_summary_inner_quantile['runtime_minutes'].mean().round(2)
```

```
Out[100]: 84.81
```

```
In [101]: 1 # Consider median  
2 runtime_summary_inner_quantile['runtime_minutes'].median().round(2)
```

```
Out[101]: 87.0
```

→ It is good that both values are about the same. Missing values will be filled with the median of 87 minutes because it is a better option as it is less sensitive to outliers.

```
In [102]: 1 # Make a copy to convert from a view to a data frame  
2 runtime_summary_inner_quantile = runtime_summary_inner_quantile.copy()  
3  
4 # Fill missing values with median  
5 runtime_summary_inner_quantile['runtime_minutes'] = runtime_summary_inner_quantile['runtime_minutes'].
```

Which runtime range receives the highest mean revenue?

Runtime will be classified into different ranges based on the output from the describe method as shown here:

| Runtime(in minutes) | Range |
|---------------------|-----------|
| 45-65 | Short |
| 65-85 | Medium |
| 85-105 | Long |
| 105-125 | Very Long |

In [103]:

```

1 # Use the Pandas cut method to map the Runtime into the above ranges
2
3 # Define cutoff for each bin
4 runtime_bin_edges = [45, 65, 85, 105, 125]
5
6 # Name each bin
7 runtime_bin_names = ['Short (45-65)', 'Medium (65-85)', 'Long (85-105)', 'Very Long (105-125)']
8
9 # Make column for different ranges into which the values of the 'runtime_minutes' column are divided
10 runtime_summary_inner_quantile['runtime_range'] = pd.cut(runtime_summary_inner_quantile.runtime_minutes,
11
12 #Display df
13 print(len(runtime_summary_inner_quantile))
14 display(runtime_summary_inner_quantile.head())
15

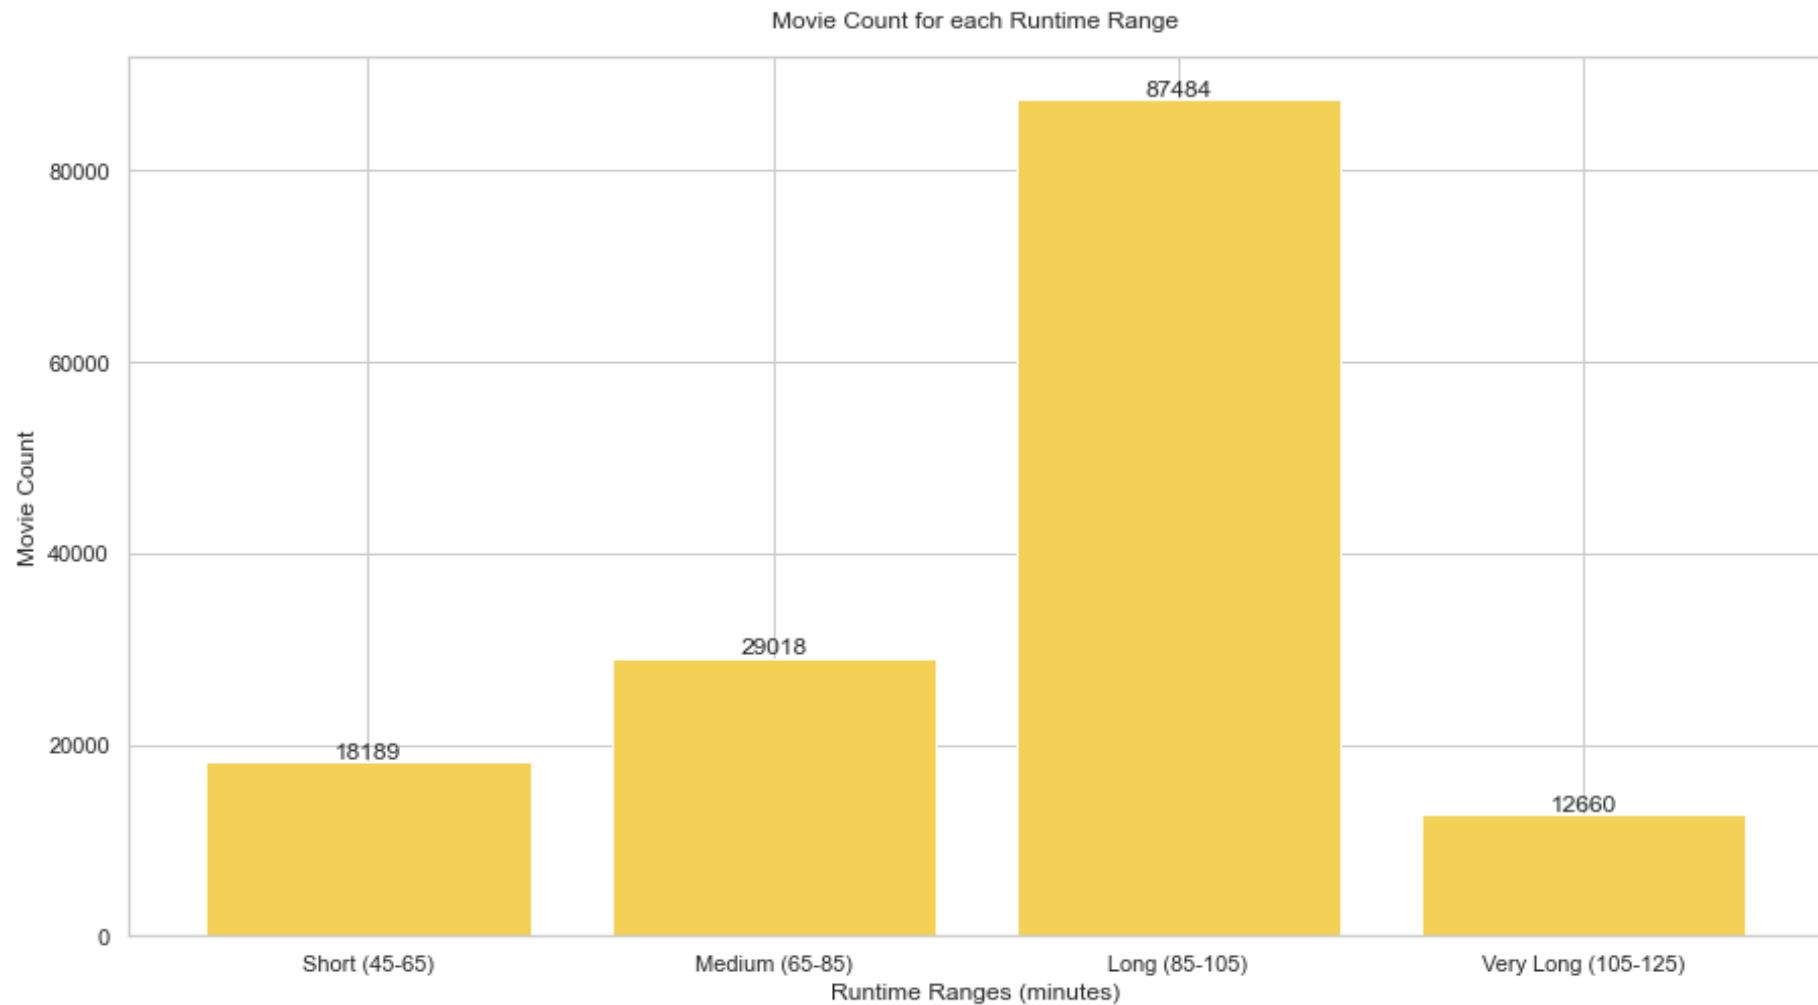
```

150377

| | title | runtime_minutes | cpi_avg_domestic_gross | genres | release_year | runtime_range |
|---|---------------------------------|-----------------|------------------------|----------------------|--------------|---------------------|
| 1 | one day before the rainy season | 114.0 | NaN | Biography,Drama | 2019.0 | Very Long (105-125) |
| 2 | sabse bada sukh | 87.0 | NaN | Comedy,Drama | 2018.0 | Long (85-105) |
| 3 | the wandering soap opera | 80.0 | NaN | Comedy,Drama,Fantasy | 2017.0 | Medium (65-85) |
| 4 | a thin life | 75.0 | NaN | Comedy | 2018.0 | Medium (65-85) |
| 5 | bigfoot | 87.0 | NaN | Horror,Thriller | 2017.0 | Long (85-105) |

In [104]:

```
1 def runtime_bin_visual(df):
2     """
3         Creates a bar plot of movie counts for each runtime range based on the 'runtime_range' column.
4
5     Input:
6         df : Pandas dataframe containing information about movies with a 'runtime_range' column
7
8     Output:
9         None (Displays a matplotlib figure)
10    """
11    # Create a Counter of runtime ranges from the available data
12    bins_counter = Counter(df['runtime_range'])
13
14    # Define all bins to display
15    all_bins = ['Short (45-65)', 'Medium (65-85)', 'Long (85-105)', 'Very Long (105-125)']
16
17    # Create counts for all bins, filling missing bins with zero counts
18    counts_per_bin = {bins: bins_counter.get(bins, 0) for bins in all_bins}
19
20    # Plot the count of different runtime ranges
21    plt.figure(figsize=(15, 8))
22    ax = plt.bar(counts_per_bin.keys(), counts_per_bin.values(), color="#F4D054")
23
24    plt.ylabel('Movie Count')
25    plt.xlabel('Runtime Ranges (minutes)')
26    plt.title('Movie Count for each Runtime Range', y=1.02)
27
28    # Show the counts on top of the bars
29    for i, v in enumerate(counts_per_bin.values()):
30        plt.text(i, v + 0.1, str(v), ha='center', va='bottom')
31
32
33    plt.show()
34
35    # View dataframe
36    runtime_bin_visual(runtime_summary_inner_quantile)
37
```



→ Let's see if this pattern follows for the top grossing movies.

In [105]:

```
1 # Inner merge 'top_profit' with 'runtime_summary_inner_quantile' on 'title' and 'year'
2 top_profit_runtime = pd.merge(top_profit, runtime_summary_inner_quantile[['title', 'release_year']], le
3
4 # Display the resulting DataFrame
5 print(len((top_profit_runtime)))
6 display((top_profit_runtime).head())
7
```

93

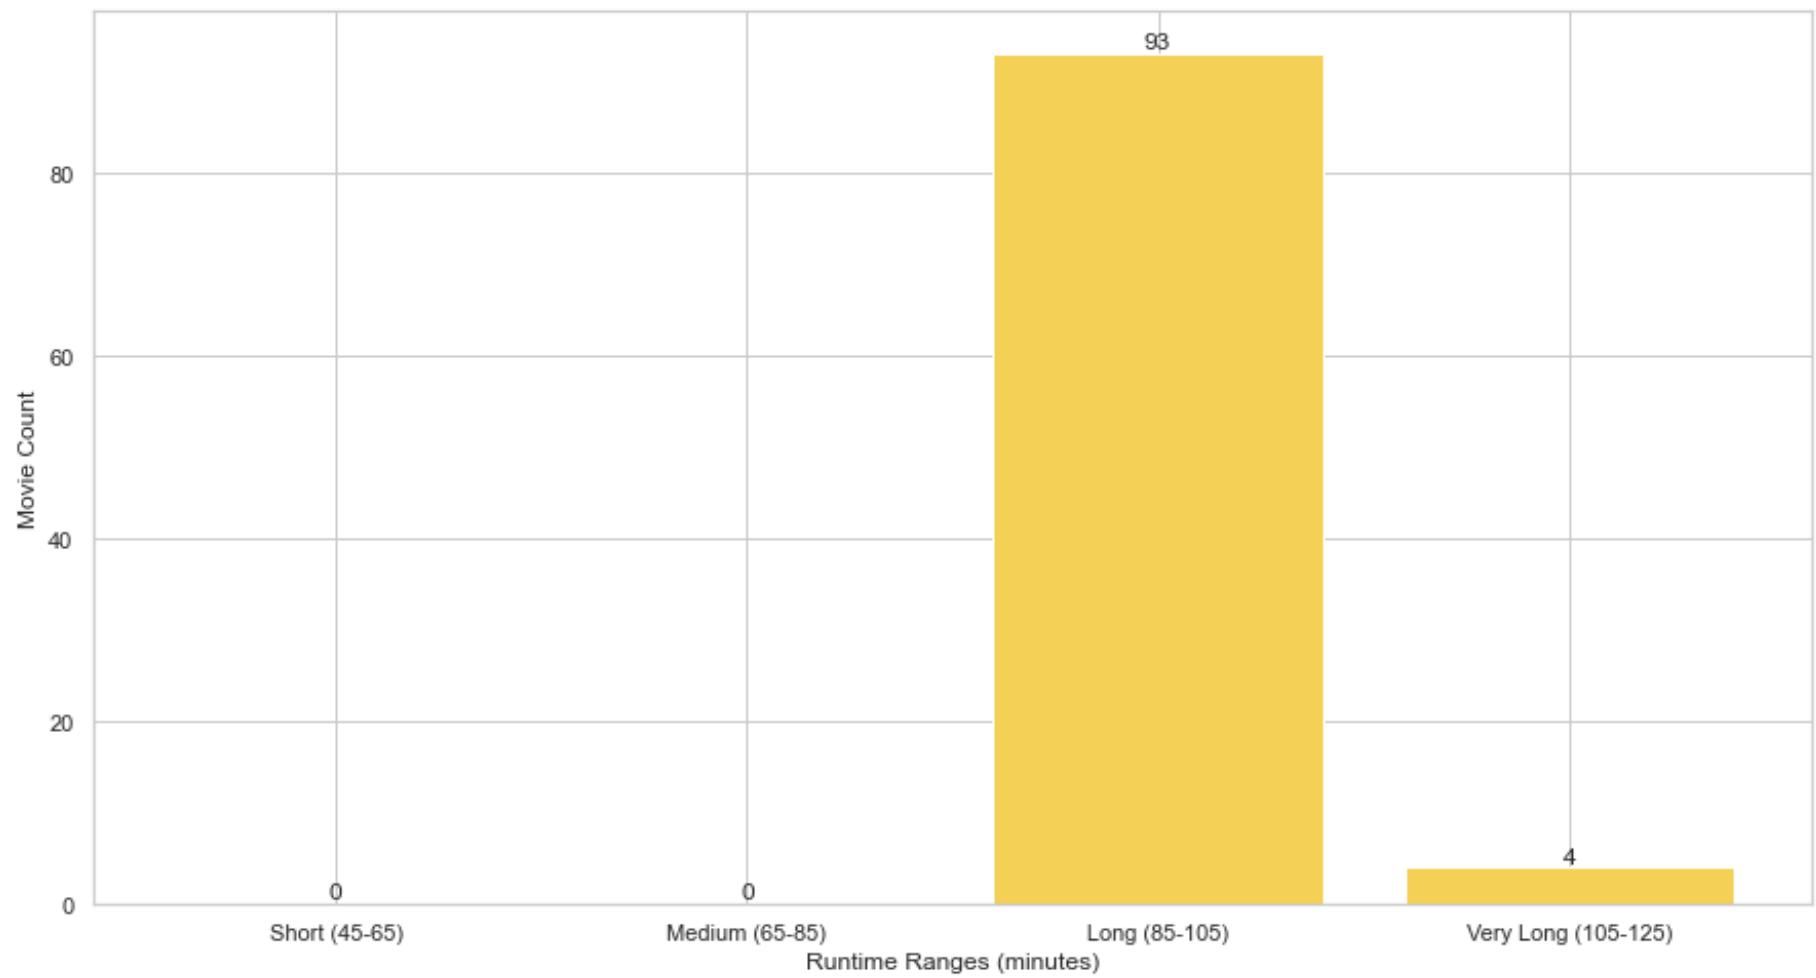
| | title | foreign_gross | year | release_date | production_budget | worldwide_gross | avg Domestic gross | cpi_worldwide_gross | cpi_avg Domestic |
|---|--|---------------|------|--------------|-------------------|-----------------|--------------------|---------------------|------------------|
| 0 | gone with the wind | NaN | 1939 | 1939-12-15 | 3900000.0 | 390525192.0 | 198680470.0 | 8.560706e+09 | 4.35521 |
| 1 | snow white and the seven dwarfs | NaN | 1937 | 1937-12-21 | 1488000.0 | 184925486.0 | 184925486.0 | 3.912998e+09 | 3.91299 |
| 2 | star wars ep. iv: a new hope | NaN | 1977 | 1977-05-25 | 11000000.0 | 786598007.0 | 460998007.0 | 3.955082e+09 | 2.31791 |
| 3 | bambi | NaN | 1942 | 1942-08-13 | 858000.0 | 268000000.0 | 102797000.0 | 5.009824e+09 | 1.92162 |
| 4 | pinocchio | NaN | 1940 | 1940-02-09 | 2289247.0 | 84300000.0 | 84300000.0 | 1.834741e+09 | 1.83474 |

→ The top_profit_runtime is only showing 93/100 entries because some of the entries do not contain runtimes.

In [106]:

```
1 # Set top_profit_runtime as top 100 movies
2 top_profit_runtime = runtime_summary_inner_quantile.sort_values('cpi_avg_domestic_gross', ascending = True)
3
4 def runtime_bin_visual_tg(df):
5     """
6         Creates a bar plot of movie counts for each runtime range based on the 'runtime_range' column.
7
8         Input:
9             df : Pandas dataframe containing information about top-profit movies with a 'runtime_range' column
10
11        Output:
12            None (Displays a matplotlib figure)
13    """
14
15    # Create a Counter of runtime ranges from the available data
16    bins_counter = Counter(df['runtime_range'])
17
18    # Define all bins to display
19    all_bins = ['Short (45-65)', 'Medium (65-85)', 'Long (85-105)', 'Very Long (105-125)']
20
21    # Create counts for all bins, filling missing bins with zero counts
22    counts_per_bin = {bins: bins_counter.get(bins, 0) for bins in all_bins}
23
24    # Plot the count of different runtime ranges
25    plt.figure(figsize=(15, 8))
26    ax = plt.bar(counts_per_bin.keys(), counts_per_bin.values(), color="#F4D054")
27
28    plt.ylabel('Movie Count')
29    plt.xlabel('Runtime Ranges (minutes)')
30    plt.title('Top Profit Movie Runtime Range Count', y=1.02)
31
32    # Show the counts on top of the bars
33    for i, v in enumerate(counts_per_bin.values()):
34        plt.text(i, v + 0.1, str(v), ha='center', va='bottom')
35
36    plt.show()
37
38 # View dataframe
39 runtime_bin_visual_tg(top_profit_runtime)
```

Top Profit Movie Runtime Range Count



In [107]:

```
1 # Find out which runtime range receives the highest mean domestic gross
2 runtime_stats = runtime_summary_inner_quantile.groupby('runtime_range')['cpi_avg_domestic_gross'].mean()
3 print('Total number of results:', len(runtime_stats))
4 display(runtime_stats)
```

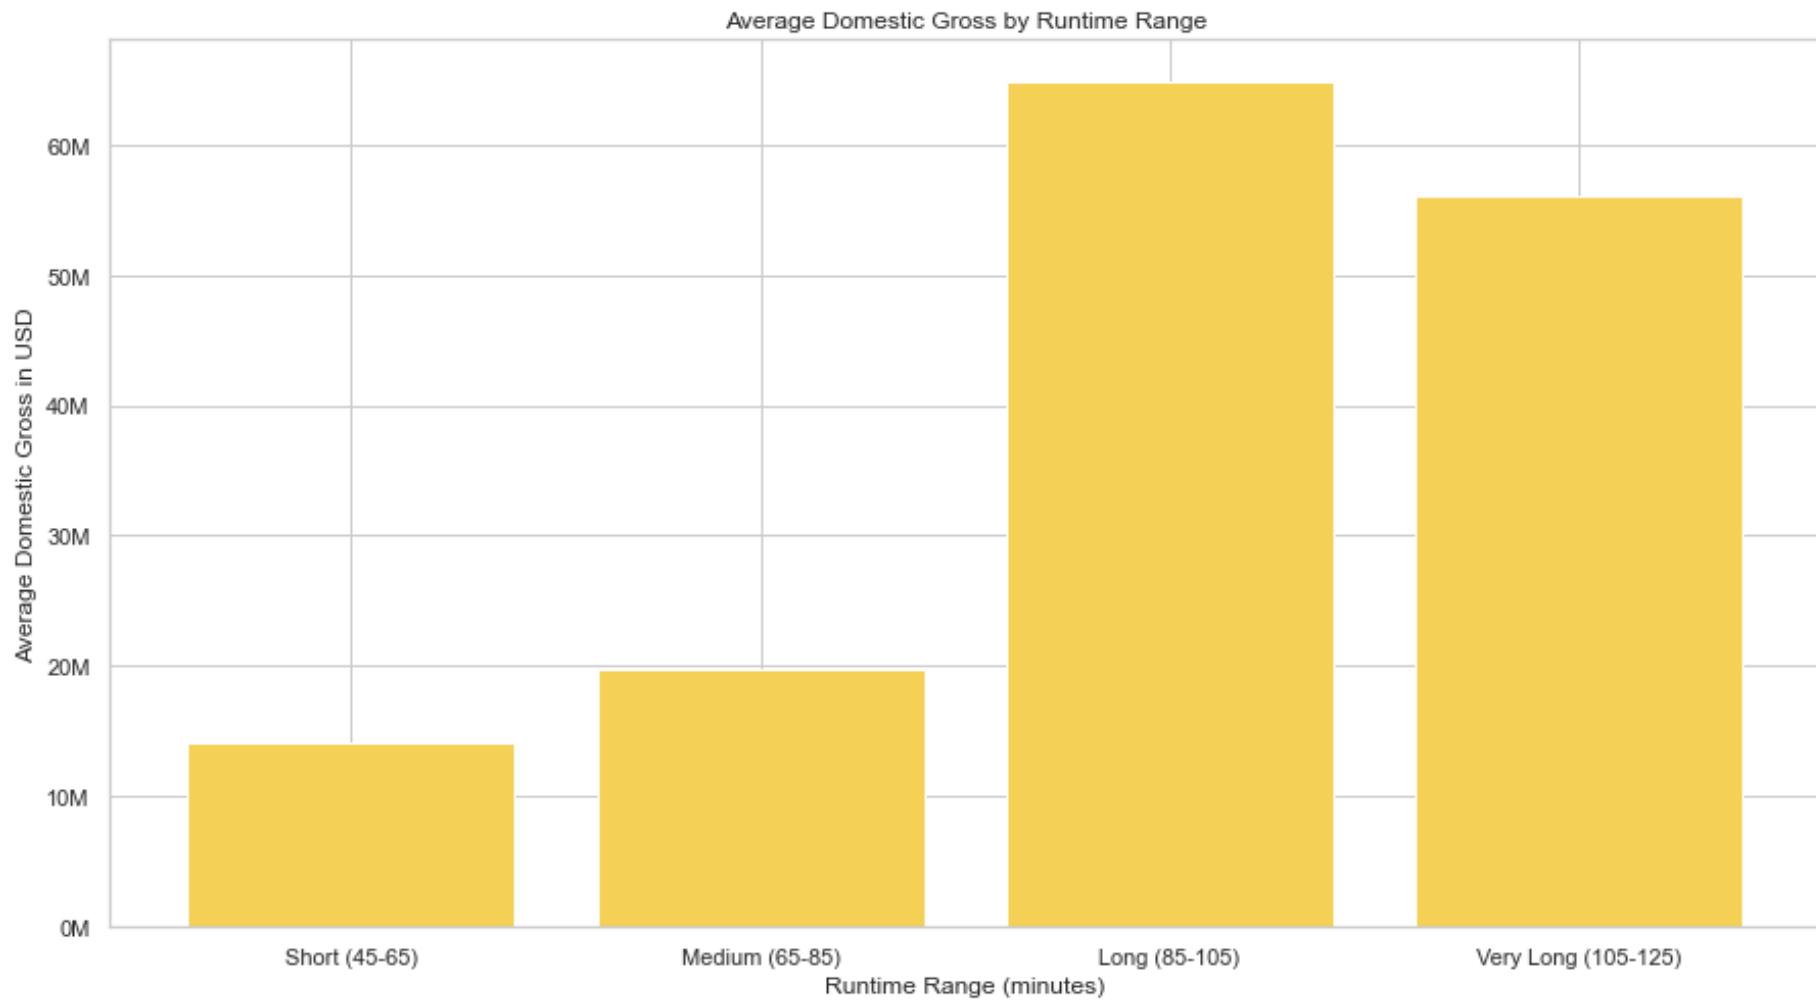
Total number of results: 4

```
runtime_range
Short (45–65)      1.412140e+07
Medium (65–85)     1.978483e+07
Long (85–105)      6.479307e+07
Very Long (105–125) 5.601178e+07
Name: cpi_avg_domestic_gross, dtype: float64
```

These bin categories can be altered to suit the stakeholders budget because longer movies may incur larger production budgets. These initial bin category settings are a conservative estimate with respect to what is considered a long movie. This analysis, prioritizes bins with an equal width of 20 minute windows. For example, a deeper dive may include eight, 10 minute wide bins.

In [108]:

```
1 def millions_formatter(x, pos):
2     """
3         Custom formatter for formatting y-axis values in millions.
4
5     Input:
6         - x (float): Value to be formatted
7         - pos (int): Position of the tick on the axis
8
9     Output:
10        - str: Formatted string representing the value in millions
11    """
12    return f'{x / 1e6:.0f}M'
13
14 # Recreate the DataFrame from the Series to use it in Matplotlib's bar plot
15 runtime_stats = runtime_stats.reset_index()
16
17 # Create the plot
18 plt.figure(figsize=(15, 8))
19 plt.bar(runtime_stats['runtime_range'], runtime_stats['cpi_avg_domestic_gross'], color="#F4D054")
20 plt.xlabel('Runtime Range (minutes)')
21 plt.ylabel('Average Domestic Gross in USD')
22 plt.title('Average Domestic Gross by Runtime Range')
23
24 # Apply the millions formatter to the y-axis
25 formatter = FuncFormatter(millions_formatter)
26 plt.gca().yaxis.set_major_formatter(formatter)
27
28 # Show the plot
29 plt.show()
30
```



This shows a clearly consistent result where long and very long movies gross higher than short and medium length movies. Again, let's now compare this to the top 100 grossing movies.

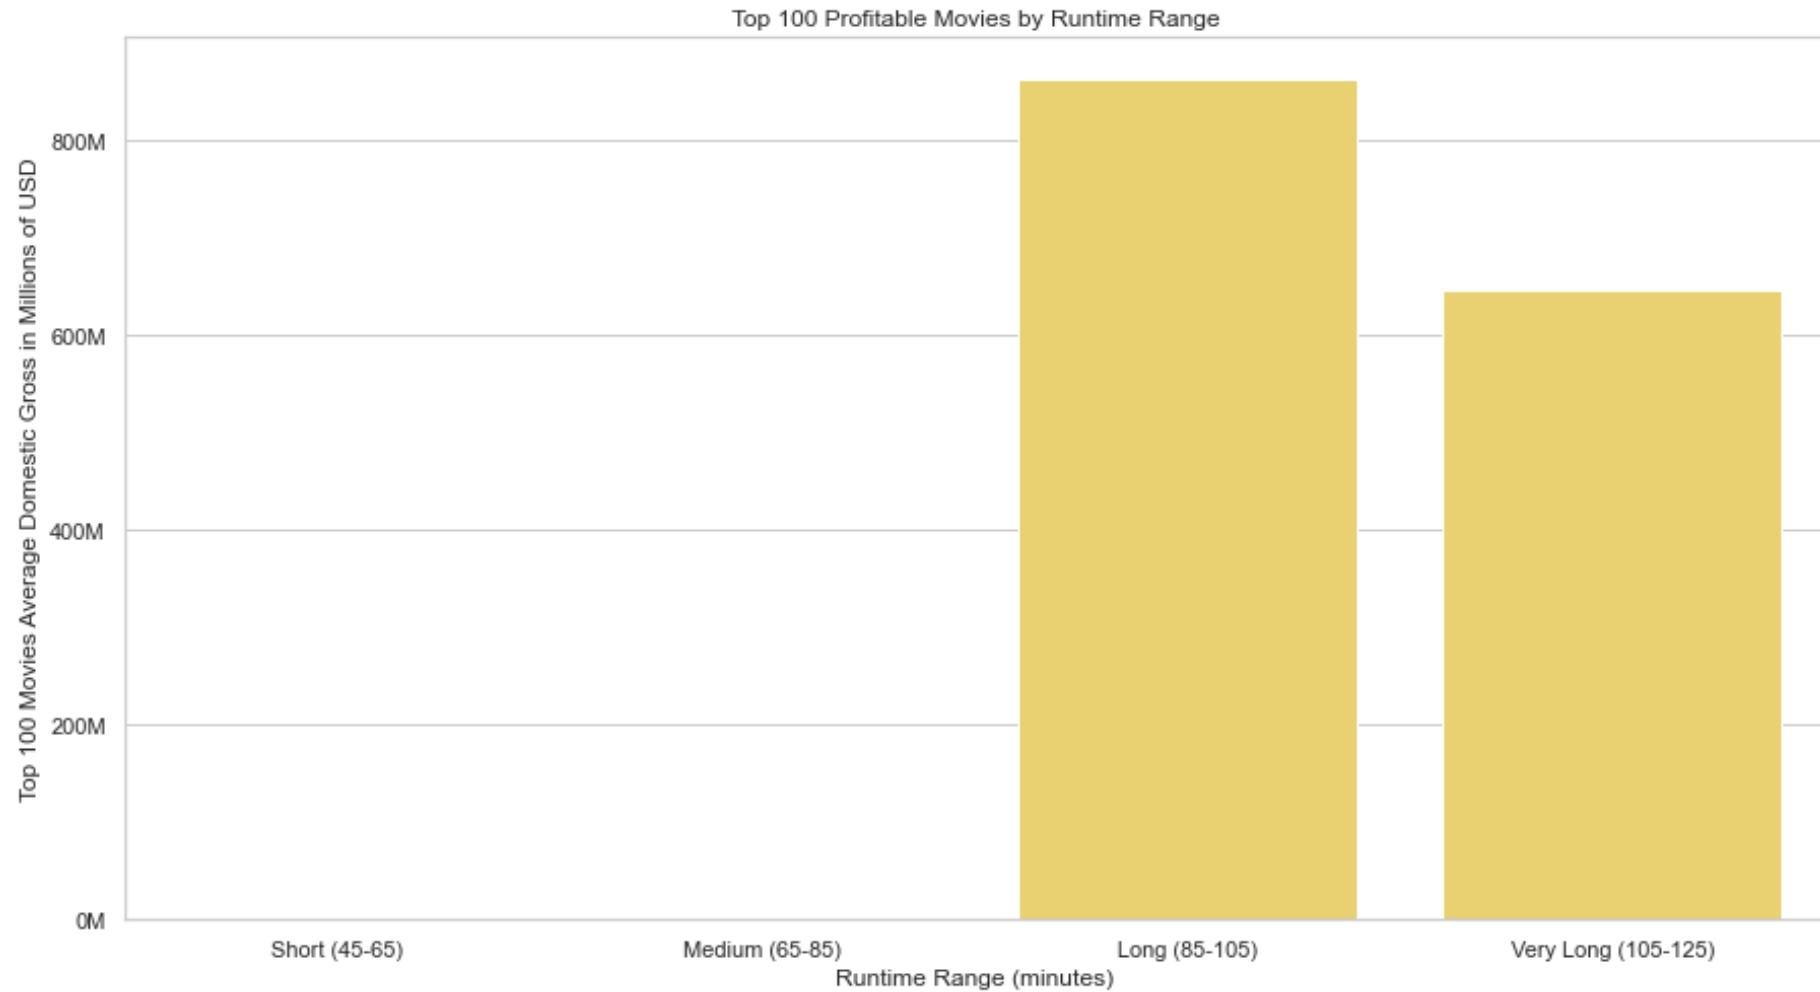
In [109]:

```
1 # Find out which runtime range receives the highest mean gross
2 runtime_stats2 = top_profit_runtime.groupby('runtime_range')['cpi_avg_domestic_gross'].mean()
3 print('Total number of results:', len(runtime_stats2))
4 display(runtime_stats2)
```

Total number of results: 4

```
runtime_range
Short (45–65)           NaN
Medium (65–85)          NaN
Long (85–105)           8.625266e+08
Very Long (105–125)     6.470557e+08
Name: cpi_avg_domestic_gross, dtype: float64
```

```
In [110]: create the DataFrame from the Series to use it in a barplot
imdb_stats2 = runtime_stats2.reset_index()
    3
    4    # Create the plot
    5    figure(figsize=(15, 8))
    6    sns.barplot(x='runtime_range', y='cpi_avg_domestic_gross', estimator= sum, ci=None, color='#FDDC5C', data=
    7    xlabel('Runtime Range (minutes)')
    8    ylabel('Top 100 Movies Average Domestic Gross in Millions of USD')
    9    title('Top 100 Profitable Movies by Runtime Range')
   10
   11    # Define a custom formatter function to display values in millions
   12    millions_formatter(x, pos):
   13        '13
   14        Custom formatter for formatting y-axis values in millions.
   15
   16    Input:
   17        x: Value to be formatted
   18        pos: Position of the tick on the axis
   19
   20    Output:
   21        formatted_string : Formatted string representing the value in millions
   22
   23        return f'{x / 1e6:.0f}M'
   24
   25    # Apply the millions formatter to the y-axis
   26    formatter = FuncFormatter(millions_formatter)
   27    gca().yaxis.set_major_formatter(formatter)
   28
   29    # Show the plot
   30    show()
```



As expected, the data supports that the top 100 grossing movies in the long to very long range are likely to have a higher average domestic gross, there for a higher return on investment.

In [111]:

```
1 # Close the database connection
2 conn.close()
```

Summary

This project aims to identify factors influencing box office revenue for movies by analyzing datasets from Box Office Mojo, IMDB, Rotten Tomatoes, The Movie DB, and The Numbers. The analysis focuses on genres, release months, runtime, and their impact on Return on Investment (ROI).

Recommendations:

The data shows that The Film Company should consider the following to maximize profits:

1. Release movies in May/June or November/December.
2. Release action or adventure.
3. Produce a movie with a runtime between 85-125 minutes.
4. Plan for a budget of 44-136 million.

Profitability Trends:

- A strong, positive correlation exists between profit and average domestic gross, indicating that top-grossing films tend to be the most profitable.
 - Caution: this is a simplified approach that does not account for other costs.
- Profit margins for the top 100 grossing movies vary, with a median profit margin of approximately 88%.
- The median profit for a top 100 grossing movie is 500 million.
- The relationship between budget and adjusted average domestic gross of the top 100 grossing movies, does not show a clear correlation. It would seem that movies have seen success with small budgets and others despite robust budgets, failed to perform.
- Budgets between 44 and 136 million seem to be reasonable for a top 100 grossing movie.

Genre Impact:

- Adventure, Action, and Comedy genres emerge as financially successful, while certain genres like Talk-Show, Short, Reality-TV, Adult, and Game-Show show zero CPI-adjusted average domestic gross, possibly due to missing financial data.
- Adventure, Action, and Comedy account for almost 64% of all adjusted domestic profits

Release Month Influence:

1. Movies released in May, June, July, and November/December tend to be more profitable, emphasizing the importance of strategic release timing.
2. Caution is advised, recognizing external factors like the economy, distribution strategy, advertising budget, and movie's enduring appeal.

Runtime Considerations:

1. After addressing outliers, the runtime data appears normally distributed, with a mean of approximately 87 minutes.

Runtime and Revenue:

1. Classifying movies into runtime ranges (Short, Medium, Long, Very Long) reveals a consistent trend where long and very long movies tend to have higher average domestic gross.

Limitations and Further Considerations:

1. Potential limitations include data gaps, advertising influence, streaming releases, and limited theater releases.
2. Adapting bin categories for runtime based on stakeholder budget considerations is suggested for a more tailored analysis.

Next Steps:

While providing valuable insights, the analysis acknowledges the need for further investigation to address nuances and ensure a comprehensive understanding of movie dynamics influencing financial success.

1. In-Depth Genre Analysis:

Conduct a detailed investigation into genres with zero CPI-adjusted average domestic gross. Verify the data sources and integrity, explore reasons for the absence of financial information, and determine if these genres are genuinely unprofitable or if there are data gaps.

2. Mixed-Genre Film Exploration:

Devote a segment of the analysis to understanding the influence of mixed-genre films on financial performance. This could involve categorizing movies with multiple genres and examining how these combinations contribute to box office revenue. Insights gained could refine genre-specific recommendations.

3. External Factors Impacting Release Timing:

Extend the analysis to consider external factors influencing release timing, beyond month-based trends. Factors like major economic events, holidays, or cultural phenomena may impact movie profitability. Identifying and understanding these factors can provide more nuanced recommendations for optimal release schedules.

4. Consumer Preferences and Genre Evolution:

Investigate evolving consumer preferences in movie genres over time. Analyze historical trends to identify shifts in audience interests. Understanding how genres have evolved can inform strategic decisions about genre selection/combination and potential shifts in the future.

5. Comprehensive ROI Forecasting Model:

Develop a comprehensive model that incorporates specific factors prioritized by the film company based on their staff strengths and goals. This could involve collaboration with stakeholders to identify key performance indicators (KPIs) and create a forecasting model tailored to the company's unique priorities.

These next steps aim to address identified inconsistencies, enhance the robustness of the analysis, and provide more actionable insights for the film company's decision-making processes.

References

1. <https://pypi.org/project/cpi/#description> (<https://pypi.org/project/cpi/#description>) (CPI)
2. <https://sqlite.org/forum/thread/2ca63507ad> (<https://sqlite.org/forum/thread/2ca63507ad>) (PRAGMA table)
3. <https://stackoverflow.com/questions/24258878/how-to-split-comma-separated-values/32051164#32051164> (<https://stackoverflow.com/questions/24258878/how-to-split-comma-separated-values/32051164#32051164>) (SQL split comma separated values)
4. <https://www.youtube.com/watch?v=Ohj-CqALrwk> (<https://www.youtube.com/watch?v=Ohj-CqALrwk>) (SQL database backend)
5. <https://www.youtube.com/watch?v=E-BEOD0EPDA> (<https://www.youtube.com/watch?v=E-BEOD0EPDA>) (dataframe -> SQL migration)
6. <https://www.sqlitetutorial.net/sqlite-import-csv/> (<https://www.sqlitetutorial.net/sqlite-import-csv/>) (import CSV to SQL)
7. <https://www.youtube.com/watch?v=YyUknBHCZB8> (<https://www.youtube.com/watch?v=YyUknBHCZB8>) (pandas -> SQL)
8. <https://stackoverflow.com/questions/13643558/should-glob-glob-be-preferred-over-os-listdir-or-the-other-way-around> (<https://stackoverflow.com/questions/13643558/should-glob-glob-be-preferred-over-os-listdir-or-the-other-way-around>) (glob)
9. https://help.hightbond.com/helpdocs/analytics/141/scripting-guide/en-us/Content/lang_ref/functions/r_exclude.htm (https://help.hightbond.com/helpdocs/analytics/141/scripting-guide/en-us/Content/lang_ref/functions/r_exclude.htm) (exclude function)
10. <https://www.geeksforgeeks.org/difference-between-list-and-dictionary-in-python/> (<https://www.geeksforgeeks.org/difference-between-list-and-dictionary-in-python/>) (list and dictionary background)
11. <https://www.youtube.com/watch?v=tATFQUx0Zx0> (<https://www.youtube.com/watch?v=tATFQUx0Zx0>) (more glob)
12. <https://developer.imdb.com/non-commercial-datasets/> (<https://developer.imdb.com/non-commercial-datasets/>) (IMDb data)
13. <https://www.youtube.com/watch?v=THHwJcKfGLQ> (<https://www.youtube.com/watch?v=THHwJcKfGLQ>) (tsv -> csv)
14. <https://www.pythontutorial.net/python-basics/python-scientific-notation/> (<https://www.pythontutorial.net/python-basics/python-scientific-notation/>) (number to words converter)
15. <https://www.youtube.com/watch?v=inMGMGhYU3uU> (<https://www.youtube.com/watch?v=inMGMGhYU3uU>) (reset index)
16. <https://www.youtube.com/watch?v=xOoCd5VBzXw> (<https://www.youtube.com/watch?v=xOoCd5VBzXw>) (project set up)
17. <https://www.themoviedb.org/talk/5daf6eb0ae36680011d7e6ee> (<https://www.themoviedb.org/talk/5daf6eb0ae36680011d7e6ee>) (movie database support)
18. <https://learnsql.com/blog/learn-and-practice-sql-joins/> (<https://learnsql.com/blog/learn-and-practice-sql-joins/>) (SQL joins)
19. <https://xkcd.com/color/rgb/> (<https://xkcd.com/color/rgb/>) (xkcd color codes)

20. [\(https://github.com/kimfetti/Videos/blob/master/Seaborn/26_palettes.ipynb\) \(seaborn color codes\)](https://github.com/kimfetti/Videos/blob/master/Seaborn/26_palettes.ipynb)
21. [\(https://www.youtube.com/watch?v=Eg0NJcUWLRM\) \(all color codes\)](https://www.youtube.com/watch?v=Eg0NJcUWLRM)
22. [\(https://seaborn.pydata.org/tutorial/color_palettes.html\) \(seaborn pallets\)](https://seaborn.pydata.org/tutorial/color_palettes.html)
23. [\(https://gizmodo.com/how-much-money-does-a-movie-need-to-make-to-be-profitab-5747305\) \(movie article\)](https://gizmodo.com/how-much-money-does-a-movie-need-to-make-to-be-profitab-5747305)
24. [\(https://stackoverflow.com/questions/22676081/what-is-the-difference-between-join-and-merge-in-pandas\) \(join and merge differences PANDAS\)](https://stackoverflow.com/questions/22676081/what-is-the-difference-between-join-and-merge-in-pandas)
25. [\(https://smallbusiness.chron.com/difference-between-profit-profit-margin-1595.html\) \(profit vs profit margin\)](https://smallbusiness.chron.com/difference-between-profit-profit-margin-1595.html)
26. [\(https://learngitbranching.js.org/\) \(visualizing git\)](https://learngitbranching.js.org/)
27. [\(https://flowingdata.com/2009/11/25/9-ways-to-visualize-proportions-a-guide/\) \(visualization ideas\)](https://flowingdata.com/2009/11/25/9-ways-to-visualize-proportions-a-guide/)